


# TERRAFORM ON AWS...

## BROWNFIELD... MULTIREGION

*Just works? Worth it?*

**BOB DYLAN**

**I CAN'T  
SING.**



# GOALS OF TALK

- Many blog posts/videos on intros to TF
- Less comprehensive resources on dealing with multiple regions, brownfield implementations and refactoring
- Goal today is to briefly introduce IAC/Terraform then talk about three main related things:
  - Multi-region project structure
  - Importance of modules in relation
  - Refactoring

# WHY (I CHOSE) INFRA AS CODE?

- Time/tediousness
- Human error
- Replicability
- Repeatability
- Stability
- 'Code as documentation'

"People seldom do what they believe in.



They do what is convenient, then repent."

# WHAT IS IAC/TERRAFORM?

Framework to specify infrastructure resources via code.

Example: Delcare a VPC on AWS

```
" Press ? for help
```

```
.. (up a dir)
```

```
</snappytv-terraform/
```

```
▶ .git/
```

```
▶ base/
```

```
▼ envs/
```

```
▶ prod/
```

```
▶ rel/
```

```
▶ stage/
```

```
  .DS_Store
```

```
▼ modules/
```

```
▶ security_groups_foreign_classic/
```

```
▶ security_groups_usw2_classic/
```

```
▶ security_groups_vpc/
```

```
▼ vpc/
```

```
  main.tf
```

```
  outputs.tf
```

```
  variables.tf
```

```
.DS_Store
```

```
.gitignore
```

```
apply_all.sh*
```

```
apply_all_usw2.sh*
```

```
custom_apply.sh*
```

```
global_variables.tf
```

```
README.md
```

```
~
```

```
resource "aws_vpc" "main" {  
  cidr_block      = "${var.vpc_cidr_block}"  
  enable_dns_support = true  
  enable_dns_hostnames = true
```

```
  tags = {  
    Name = "${var.enviro}"  
    Environment = "${var.enviro}"  
    TFManaged = true  
  }  
}
```

```
resource "aws_subnet" "public" {  
  ← 12 lines: count = "${length(var.subnet_cidr_blocks_public)}" --  
}
```

```
resource "aws_subnet" "private" {  
  ← 11 lines: count = "${length(var.subnet_cidr_blocks_private)}" --  
}
```

```
resource "aws_internet_gateway" "gw" {  
  ← 7 lines: vpc_id = "${aws_vpc.main.id}"  
}
```

```
resource "aws_route_table" "rte" {  
  ← 12 lines: vpc_id = "${aws_vpc.main.id}"  
}
```

```
resource "aws_route_table_association" "a" {
```

# WHY INFRA AS CODE ON EOL SNAPPYTV?

- Inherited amazing puppet repo for SnappyTV
- All instance creation, ELB's, security groups, etc have been managed manually by SRE.
- Need to upgrading our dynamic instances to AWS C5 size
- The need to spin up 27 vpc's (9 regions \* 3 environs) systematically.
- The need to audit and segregate security groups and rules by environment (i.e. dev/staging/production)
- Seemed like reasonable tool to help with the above

# WHY TERRAFORM?

- Puppet has some modules which are IAC but... we're still on Puppet 3 and these not avail  
I.e <https://github.com/puppetlabs/puppetlabs-aws>
- Provider agnostic - TF works with AWS (our case) but also GCP, and probably 100+ others.
  - Investment for company (perhaps), engineer (certainly)
- Seemed like the strongest and most used if IAC frameworks
- Seemed to be good docs and community



# CHALLENGES

- Terraform 'provider' which is the interface to the external service (i.e. AWS, GCP), in the case of AWS requires a 'region' field.

Ie: `provider "aws" { region = "ap-northeast-1" }`

- SnappyTV in 9 AWS regions
- 3 environments per region: production, qa, development
- OMG... 27 individual TF state files!



# STATE FILE(S) AND WHERE???

- Local state?, git?, remote state
- Terraform remote state S3 backend (others avail)
  - Locking so only one dev can apply one tf repo at a time
  - Hands-off - once set up TF automatically handles reads and writes

# GETTING STARTED WITH TF AND SCARY THINGS

- Installing from TF site and docs straightforward
- Is quick to get to 'hello world' with a local statefile
- Most frightening thing is letting TF run the first time against a major brownfield infrastructure
- `terraform destroy` - destroys all managed infra!
- In the end TF can not destroy (at least in AWS) things that are dependent



# OPTIONS TO MANAGE MULTIPLE (AWS) REGIONS/ENVS:

## Terraform Workspaces

- Allow you to apply same code to different environment or 'place'
  - Non-starter for SnappyTV

## Roll Your Own

- Manually treat each region/env combination as it's own TF project... all within a tree with some shared resources
  -

```

.. (up a dir)
/Users/akaruna/workspace/snappytv-terraform/
└─ .git/
└─ base/
    └─ .terraform/
        global_variables.tf -> /Users/akaruna/workspace/snappytv-terraform/global_variables.tf
        s3.tf
        terraform.tfstate
        terraform.tfstate.backup
        variables.tf
└─ envs/
    └─ prod/
        └─ ap-northeast-1/
            ap-southeast-1/
            ap-southeast-2/
            eu-central-1/
            eu-west-1/
            sa-east-1/
            us-east-1/
            us-west-1/
            us-west-2/
            └─ .terraform/
                global_variables.tf -> /Users/akaruna/workspace/snappytv-terraform/global_variables.tf
                security_groups.tf
                security_groups_stv_group_puppet.tf
                security_groups_stv_rds.tf
                variables.tf
                versions.tf
                vpc.tf
            .DS_Store
        rel/
        stage/
        .DS_Store
└─ modules/
    security_groups_foreign_classic/
    security_groups_usw2_classic/
    security_groups_vpc/
    vpc/
.DS_Store
.gitignore
apply_all.sh*
apply_all_usw2.sh*
custom_apply.sh*
global_variables.tf
README.md
.global_variables.tf.swp

```

# DIR/PROJECT STRUCTURE

```

└─ base/
└─ envs/
    └─ prod/
        └─ ap-northeast-1/
            └─ .terraform/
                global_variables.tf -> /Users/akaruna/workspace/snappytv-terraform/global_variables.tf
                security_groups.tf
                variables.tf
                versions.tf
                vpc.tf
            ap-southeast-1/
            ap-southeast-2/
            eu-central-1/
            eu-west-1/
            sa-east-1/
            us-east-1/
            us-west-1/
            us-west-2/
            .DS_Store
        rel/
        stage/
        .DS_Store
    modules/

```

# ISN'T RUNNING TF 27x TO APPLY CHANGE A PITA???

Initially but... this stuck and served my purposes well...

```
# /bin/bash

# this will open a new terminal window for each env * region and run `terraform apply`
for env in stage rel prod
do
  for region in ap-northeast-1 ap-southeast-1 ap-southeast-2 eu-central-1 eu-west-1 sa-east-1 us-east-1 us-west-1 us-west-2
  do
    echo "Running $env $region"
    osascript -e "tell application \"Terminal\" to do script \"cd ~/workspace/snappytv-terraform/envs/$env/$region && terraform apply\""
  done
done
```

- Reviewing and accepting 27x takes a minute but...
- You \*want to\* always review your diffs before changes
- Easier to review apples to apples vs one huge diff (even if were possible)
- Compared to manually... human error... digging to find problems... is major time saver and priceless



2a;sort=tag... ☆

google Calendar ✂ Twitter Eng - Hell... »

snappytv ▾ Oregon ▾ Support ▾

🔧 ↺ ⚙️ ?

🔍 1 to 4 of 4 > > >

▾ Status Checks ▾	Alarm Status
✅ 2/2 checks passed	None
✅ 2/2 checks passed	None
✅ 2/2 checks passed	None
✅ 2/2 checks passed	None

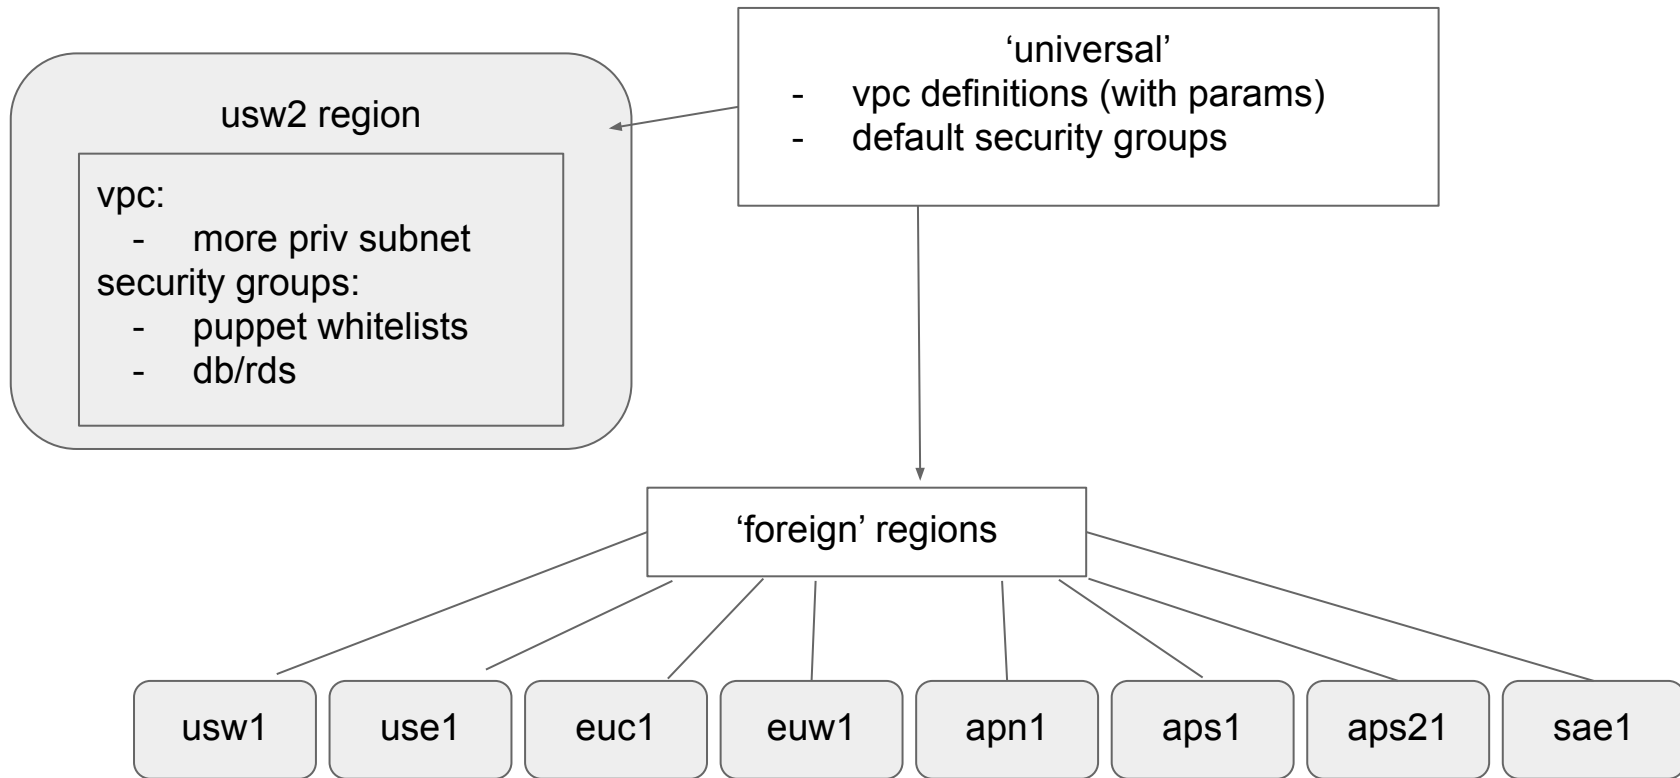
🏠 📺 📺

# STRUCTURING CODE FOR MULTIPLE 'REPOS/REGIONS'

## Challenges...

- Deriving what is actually replicated vs unique
- Getting Terraform code (HCL - Hashicorp Configuration Language) to abstract it \*maintainably by others\*
- Learning what can/can't I do with HCL?

# REPEATED VS UNIQUE SETTINGS



# MODULES: HANDLING REGION/ENV SIMILARITIES

- 'My us-west-2 region is different than my other regions'
- 'Each env uses different security groups but the base code is the same'

# 'FOREIGN' VS USW2 VPC CONFIG (ENV)

```
module "vpc" {
  source = "../../modules/vpc"

  environ = var.environ

  # us-west-1 has az's 1-3 but can not create vpc subnets in us-west-1b
  availability_zones = ["us-west-1a", "us-west-1c"]
  vpc_cidr_block      = "172.18.48.0/20"
  subnet_cidr_blocks_public = ["172.18.48.0/22", "172.18.52.0/22", "172.18.56.0/22"]
  subnet_cidr_blocks_private = ["172.18.60.0/22"]
}

~

envs/prod/us-west-1/vpc.tf
module "vpc" {
  source = "../../modules/vpc"

  environ = var.environ
  availability_zones = ["us-west-2a", "us-west-2b", "us-west-2c"]
  vpc_cidr_block      = "172.18.0.0/19"
  subnet_cidr_blocks_public = ["172.18.0.0/21", "172.18.8.0/21", "172.18.16.0/21"]
  subnet_cidr_blocks_private = ["172.18.24.0/23", "172.18.26.0/23", "172.18.28.0/23"]
}

~
~
~

envs/prod/us-west-2/vpc.tf
```

# VPC MODULE (WITH PARAMS)

```
resource "aws_vpc" "main" {
  cidr_block = "${var.vpc_cidr_block}"
}

resource "aws_subnet" "public" {
  count      = "${length(var.subnet_cidr_blocks_public)}"

  vpc_id     = "${aws_vpc.main.id}"
  cidr_block = "${element(var.subnet_cidr_blocks_public, count.index)}"
  availability_zone = "${element(var.availability_zones, count.index)}"
  map_public_ip_on_launch = true


  tags = {
    Name = "${var.envIRON} public ${count.index+1}"
    Environment = "${var.envIRON}"
    TFManaged = true
  }
}

resource "aws_subnet" "private" {
  count      = "${length(var.subnet_cidr_blocks_private)}"

  vpc_id     = "${aws_vpc.main.id}"
  cidr_block = "${element(var.subnet_cidr_blocks_private, count.index)}"
  availability_zone = "${element(var.availability_zones, count.index)}"

  tags = {
    Name = "${var.envIRON} private ${count.index+1}"
    Environment = "${var.envIRON}"
    TFManaged = true
  }
}
```





# BRINGING EXISTING SECURITY GROUPS INTO TF

- Import a group:

```
`terraform import aws_security_group.elb_sg sg-983db2a8`
```

- Must write code for what imported
  - Terraform outputs the code from the import command
- Many imports, just write a bash script



# REARRANGING/REFACTORING IN TERRAFORM CODE

- You will refactor!
- Don't be afraid of 'just starting' where you are
- After moving code, modules or renaming the output of `plan` or `apply` can be frightening... but...
- Fixing things is actually (generally) easier than you think...
- `terraform mv` is your friend
- Utilize the diff in the output and feed it into `terraform mv`
- Shell script for large moves (find tix)

# REFACTORING EXAMPLE

So with current refactor, testing to move one node, this is the delete and recreate for tf plan output:

Plan: 54 to add, 0 to change, 45 to destroy. # we are truly creating 9 new sgs, dont want to destroy and recreate the rest!

+

```
module.security_groups_usw2_classic.module.security_group_stv_auto_puppet_usw2_classic.module.security_group_stv_auto_puppet_usw2_classic_ap-northeast-1.aws_security_group.stv_auto_puppet_region[0]
```

-

```
module.security_groups_usw2_classic.module.security_groups_usw2_classic.module.security_group_stv_auto_puppet_usw2_classic_ap-northeast-1.aws_security_group.stv_auto_puppet_region[0]
```

# THE 'MV' COMMAND LOOKS LIKE...

```
`terraform state move ['- item] ['+' item]`
```

```
akaruna ~/workspace/snappytv-terraform/envs/stage/us-west-2 (master) $ terraform state mv terraform
state mv
module.security_groups_usw2_classic.module.security_groups_usw2_classic.module.security_group_stv_aut
o_puppet_usw2_classic_ap-northeast-1.aws_security_group.stv_auto_puppet_region[0]
module.security_groups_usw2_classic.module.security_group_stv_auto_puppet_usw2_classic.module.securit
y_group_stv_auto_puppet_usw2_classic_ap-northeast-1.aws_security_group.stv_auto_puppet_region[0]
```

\*\*\* bash scriptable for major moves!!! \*\*\*

# CONCLUSION


- Able to implement TF with provided docs and groups
- HCL can feel kludgy but is workable
- Thinking in terms of maintainability
- Has served SnappyTV solidly without issue for 6m+
- TF 0.12 out after implementation, good improvements
- I'd personally spend more time on:
  - Getting familiar with HCL 0.12+
  - Making more data driven (i.e. more like hiera in puppet)

# RESOURCES

This is a great post which makes clear some of the pitfalls with being a Terraform newbie and also how to structure multiple AWS environments. Much of my implementation pulls from this:

<https://charity.wtf/2016/03/30/terraform-vpc-and-why-you-want-a-tfstate-file-per-env/>

# THANK YOU!



Presentation based on work performed  
on contract at Twitter 2017-1029

[alexi@alexikaruna.com](mailto:alexi@alexikaruna.com)