

ADD0

ALL DAY DEVOPS

NOVEMBER 6, 2019

Troubleshooting Real Production Problems

Ram Lakshmanan

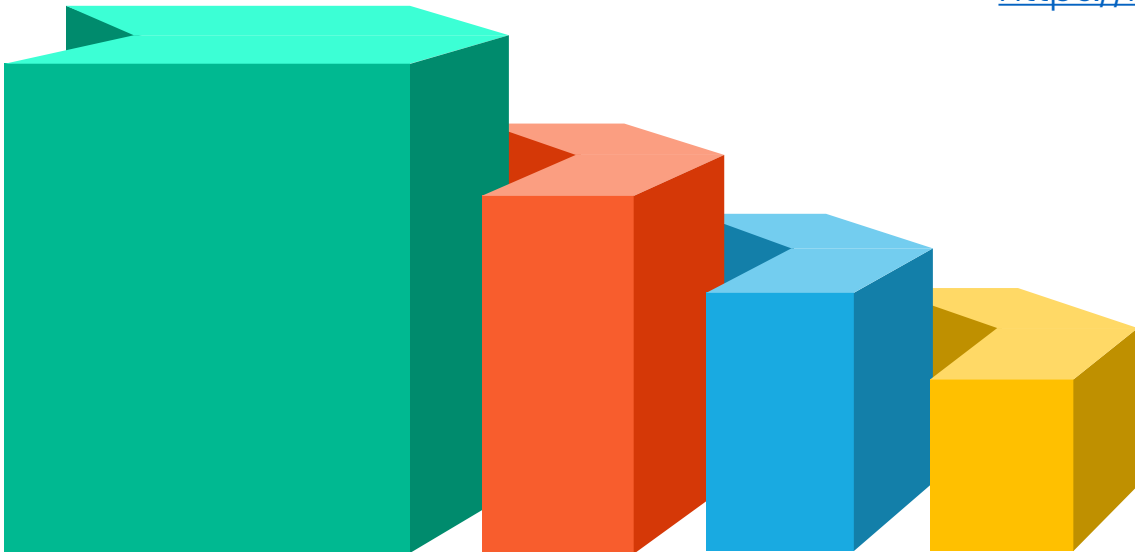
Architect: GCeasy.io, fastThread.io, HeapHero.io





Troubleshooting CPU spike

<https://blog.fastthread.io/2018/12/13/how-to-troubleshoot-cpu-problems/>



Step 1: Confirm

Don't trust anyone


```
top - 23:13:26 up 102 days, 21:09, 2 users, load average: 2.91, 2.99, 2.55
Tasks: 99 total, 1 running, 98 sleeping, 0 stopped, 0 zombie
Cpu(s):100.0%us, 0.0%sy, 0.0%ni, 0.0%id, 0.0%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 8178640k total, 3821576k used, 4357064k free, 158700k buffers
Swap: 0k total, 0k used, 0k free, 646524k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
31294	ec2-user	20	0	4415m	22m	14m	S	199.1	0.3	0:28.53	java
3108	newrelic	20	0	239m	7708	4784	S	0.7	0.1	59:09.51	nrsysmond
15153	tomcat	20	0	9132m	2.6g	17m	S	0.3	33.7	134:34.52	java
1	root	20	0	19640	2676	2344	S	0.0	0.0	0:07.55	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kthreadd
3	root	20	0	0	0	0	S	0.0	0.0	0:10.95	ksoftirqd/0
5	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/0:0H
6	root	20	0	0	0	0	S	0.0	0.0	1:27.11	kworker/u30:0
7	root	20	0	0	0	0	S	0.0	0.0	3:22.96	rcu_sched
8	root	20	0	0	0	0	S	0.0	0.0	0:00.00	rcu_bh
9	root	RT	0	0	0	0	S	0.0	0.0	0:01.96	migration/0
10	root	RT	0	0	0	0	S	0.0	0.0	0:01.44	migration/1
11	root	20	0	0	0	0	S	0.0	0.0	0:14.63	ksoftirqd/1
13	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	kworker/1:0H
14	root	20	0	0	0	0	S	0.0	0.0	0:00.00	kdevtmpfs
15	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	netns
18	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	perf
22	root	20	0	0	0	0	S	0.0	0.0	0:00.00	xenwatch
23	root	20	0	0	0	0	S	0.0	0.0	0:00.00	xenbus
134	root	20	0	0	0	0	S	0.0	0.0	0:03.71	khungtaskd
135	root	0	-20	0	0	0	S	0.0	0.0	0:00.00	writeback

'top' tool is your good friend

Step 2: Identify Threads

`top -H -p {pid}`

 ec2-user@ip-172-31-15-129:/usr/share/tomcat8/logs

```
top - 23:14:36 up 102 days, 21:10,  2 users,  load average: 2.97, 3.00, 2.58
Tasks:  15 total,   3 running, 12 sleeping,   0 stopped,   0 zombie
Cpu(s):100.0%us,  0.0%sy,  0.0%ni,  0.0%id,  0.0%wa,  0.0%hi,  0.0%si,  0.0%st
Mem:   8178640k total, 3822428k used, 4356212k free, 158700k buffers
Swap:          0k total,          0k used,          0k free, 646536k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
31306	ec2-user	20	0	4415m	22m	14m	R	69.3	0.3	0:56.53	java
31307	ec2-user	20	0	4415m	22m	14m	R	65.6	0.3	0:56.83	java
31308	ec2-user	20	0	4415m	22m	14m	R	64.0	0.3	0:55.48	java
31294	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31295	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.05	java
31296	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31297	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31298	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31299	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31300	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31301	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31302	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31303	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31304	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.00	java
31305	ec2-user	20	0	4415m	22m	14m	S	0.0	0.3	0:00.02	java

Example:
`top -H -p 31294`

Step 3: Capture thread dumps

<https://blog.fastthread.io/2016/06/06/how-to-take-thread-dumps-7-options/>

01

jstack (since Java 5)

jstack -l <pid> >
/tmp/threadDump.txt

02

kill -3

Kill -3 <pid>
Useful when only JRE is installed

03

jVisualVM

JDK tool. Now Open source.
GUI based option.

04

JMC

JDK tool. Now Open source.
GUI based option.

05

Windows (Ctrl + Break)

Helpful during development
phase

06

ThreadMXBean

Programmatic way to capture
thread dumps

07

APM Tools

Few APM Tools does provide
this support

08



Jcmd (since Java 7)

jcmd <pid> Thread.print >
/tmp/threadDump.txt

Anatomy of thread dump

2019-02-26 17:13:23

Full thread dump Java HotSpot(TM) 64-Bit Server VM (23.7-b01 mixed mode):

"InvoiceThread-A996" prio=10 tid=0x00002b7cfc6fb000 nid=0x4479 runnable [0x00002b7d17ab8000]

java.lang.Thread.State: RUNNABLE

at com.buggycompany.rt.util.ItinerarySegmentProcessor.setConnectingFlight(ItinerarySegmentProcessor.java:380)
at com.buggycompany.rt.util.ItinerarySegmentProcessor.processTripType0(ItinerarySegmentProcessor.java:366)
at com.buggycompany.rt.util.ItinerarySegmentProcessor.processItineraryByTripType(ItinerarySegmentProcessor.java:254)
at com.buggycompany.rt.util.ItinerarySegmentProcessor.templateMethod(ItinerarySegmentProcessor.java:399)
at com.buggycompany.qc.gds.InvoiceGeneratedFacade.readTicketImage(InvoiceGeneratedFacade.java:252)
at com.buggycompany.qc.gds.InvoiceGeneratedFacade.doOrchestrate(InvoiceGeneratedFacade.java:151)
at com.buggycompany.framework.gdstask.BaseGDSFacade.orchestrate(BaseGDSFacade.java:32)
at com.buggycompany.framework.gdstask.BaseGDSFacade.doWork(BaseGDSFacade.java:22)
at com.buggycompany.framework.concurrent.BuggycompanyCallable.call(buggycompanyCallable.java:80)
at java.util.concurrent.FutureTask\$Sync.innerRun(FutureTask.java:334)
at java.util.concurrent.FutureTask.run(FutureTask.java:166)
at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)
at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615)
at java.lang.Thread.run(Thread.java:722)

"Reconnection-1" prio=10 tid=0x00007f0442e10800 nid=0x112a waiting on condition [0x00007f042f719000]

java.lang.Thread.State: WAITING (parking)

at sun.misc.Unsafe.park(Native Method)

- parking to wait for <0x007b3953a98> (a java.util.concurrent.locks.AbstractQueuedSynchronizer)

at java.util.concurrent.locks.LockSupport.park(LockSupport.java:186)

at java.lang.Thread.run(Thread.java:722)

:
:
:

1 Timestamp at which thread dump was triggered

2 JVM Version info

3 Thread Details - <<details in following slides>>

1 "InvoiceThread-A996" 2 prio=10 3 tid=0x00002b7cfc6fb000 4 nid=0x4479 5 runnable [0x00002b7d17ab8000]
6 java.lang.Thread.State: RUNNABLE
 at com.buggycompany.rt.util.ItinerarySegmentProcessor.setConnectingFlight(ItinerarySegmentProcessor.java:380)
 at com.buggycompany.rt.util.ItinerarySegmentProcessor.processTripType0(ItinerarySegmentProcessor.java:366)
 at com.buggycompany.rt.util.ItinerarySegmentProcessor.processItineraryByTripType(ItinerarySegmentProcessor.java:254)
 at com.buggycompany.rt.util.ItinerarySegmentProcessor.templateMethod(ItinerarySegmentProcessor.java:399)
 at com.buggycompany.qc.gds.InvoiceGeneratedFacade.readTicketImage(InvoiceGeneratedFacade.java:252)
 at com.buggycompany.qc.gds.InvoiceGeneratedFacade.doOrchestrate(InvoiceGeneratedFacade.java:151)
7 at com.buggycompany.framework.gdstask.BaseGDSFacade.orchestrate(BaseGDSFacade.java:32)
 at com.buggycompany.framework.gdstask.BaseGDSFacade.doWork(BaseGDSFacade.java:22)
 at com.buggycompany.framework.concurrent.BuggycompanyCallable.call(buggycompanyCallable.java:80)
 at java.util.concurrent.FutureTask\$Sync.innerRun(FutureTask.java:334)
 at java.util.concurrent.FutureTask.run(FutureTask.java:166)
 at java.util.concurrent.ThreadPoolExecutor.runWorker(ThreadPoolExecutor.java:1145)
 at java.util.concurrent.ThreadPoolExecutor\$Worker.run(ThreadPoolExecutor.java:615)
 at java.lang.Thread.run(Thread.java:722)

- 1 Thread Name - *InvoiceThread-A996*
- 2 Priority - Can have values from 1 to 10
- 3 Thread Id - 0x00002b7cfc6fb000 - Unique ID assigned by JVM. It's returned by calling the Thread.getId() method.
- 4 Native Id - 0x4479 - This ID is highly platform dependent. On Linux, it's the pid of the thread. On Windows, it's simply the OS-level thread a process. On Mac OS X, it is said to be the native pthread_t value.
- 5 Address space - 0x00002b7d17ab8000 -
- 6 Thread State - RUNNABLE
- 7 Stack trace -

6 thread states

01 NEW

03 RUNNABLE

05 WAITING

`wait();`

02 TERMINATED

04 BLOCKED

Thread 2: Blocked

```
public void synchronized getData() {  
    makeDBCall();  
}
```

Thread 1: Runnable

06 TIMED_WAITING

`Thread.sleep(10);`

Step 4: Identify lines of code causing CPU spike



Thread Ids: 31306, 31307, 31308

High CPU consuming Threads Ids reported in 'top -H'.



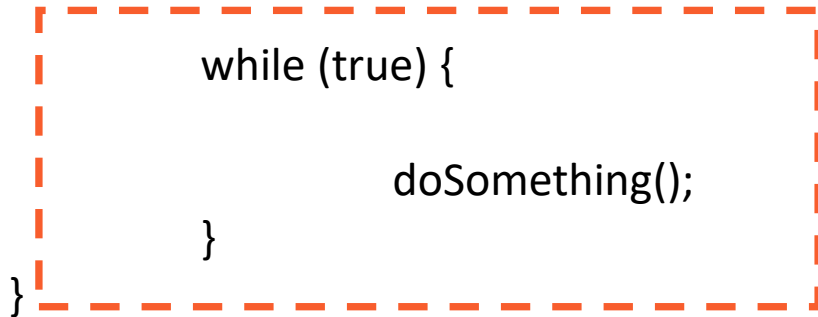
Let's look up these thread Ids in Thread dump

HexaDecimal equivalent:

- 31306 → 7a4a
- 31307 → 7a4b
- 31308 → 7a4c

Source code

```
1: package com.buggyapp.cpuspike;
2:
3: /**
4: *
5: * @author Test User
6: */
7: public class Object1 {
8:
9:     public static void execute() {
10:
11:         while (true) {
12:
13:             doSomething();
14:         }
15:     }
16:
17:     public static void doSomething() {
18:
19:     }
20: }
```



‘Free’ Thread dump analysis tools

Freely available Thread dump analysis tools

01

FastThread

<http://fastThread.io/>

02

Samurai

<http://samuraism.jp/samurai/en/index.html>

03

IBM Thread & Monitor analyzer

<https://developer.ibm.com/javasdk/tools/>

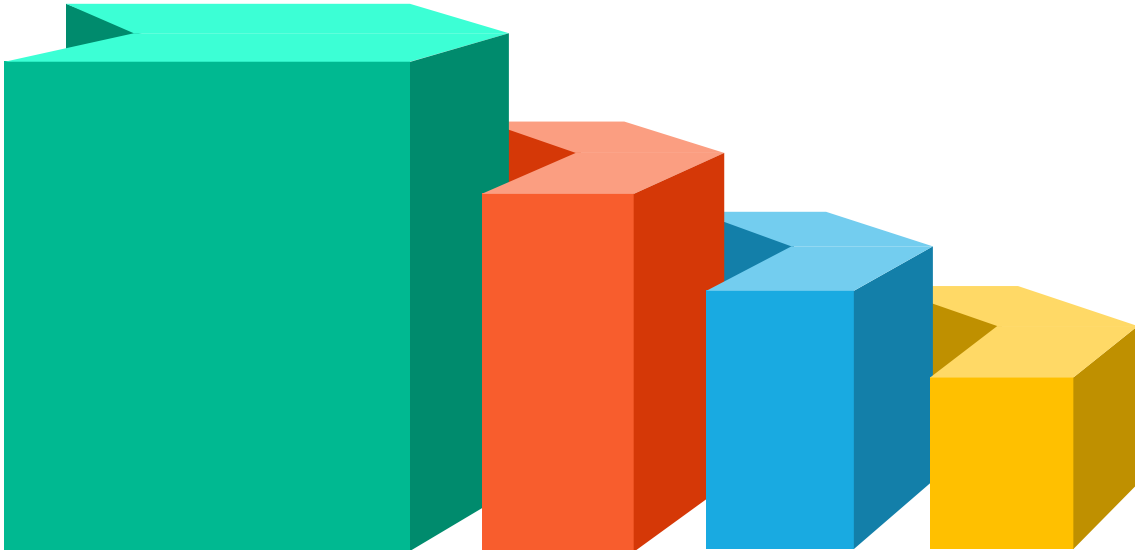
04

Visual VM

<https://visualvm.github.io/>

CPU spike in a major trading application

Troubleshooting unresponsive app

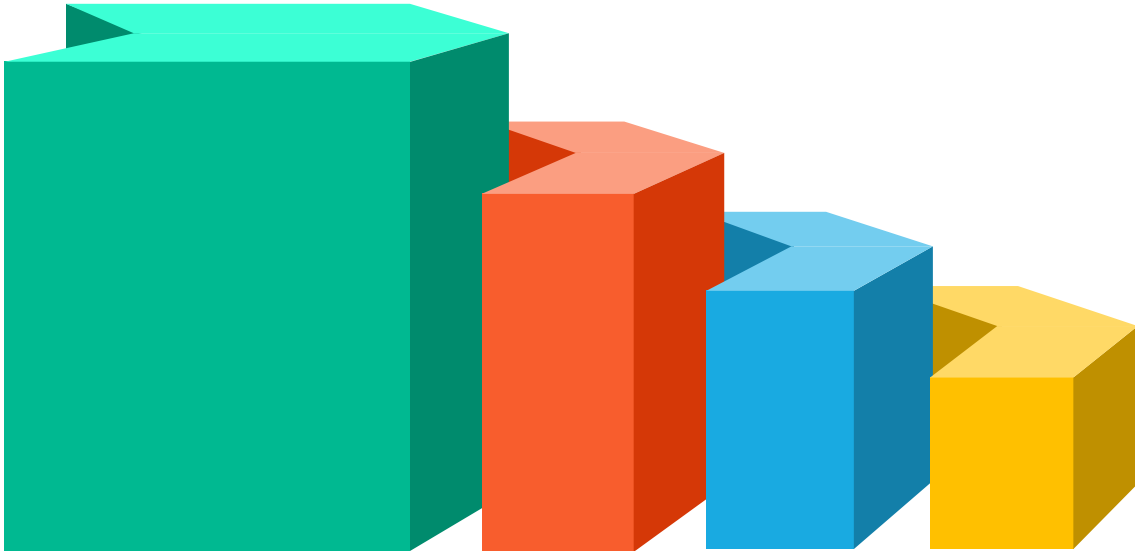


Unresponsiveness in a B2B Travel application

Process 70% of N. America oversease Leisure travel ticketing

Troubleshooting OutOfMemoryError

Unable to create new native thread



Major financial institution in N. America

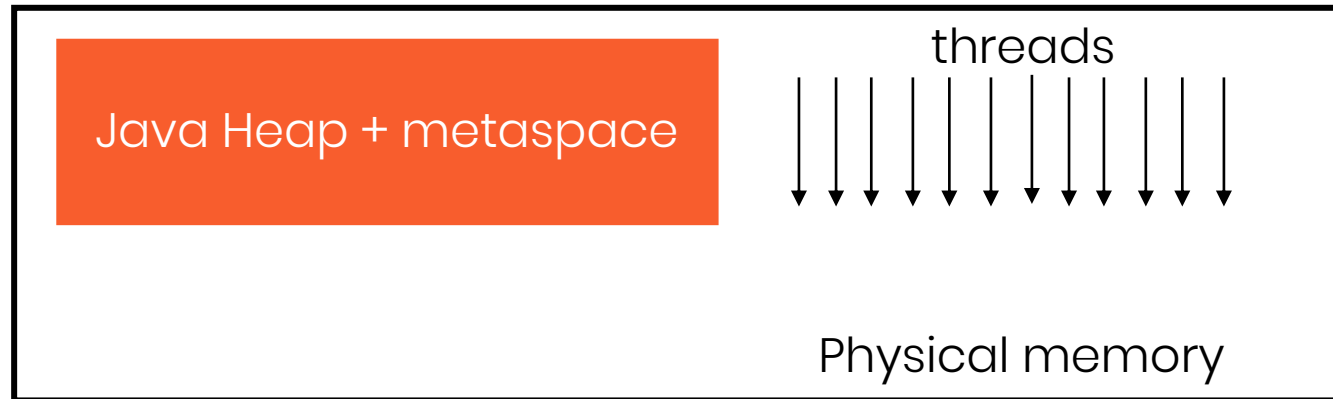


Thread dump troubleshooting pattern: RSI



<https://map.tinyurl.com/yxho6lan>

OOM: Unable to create new native thread

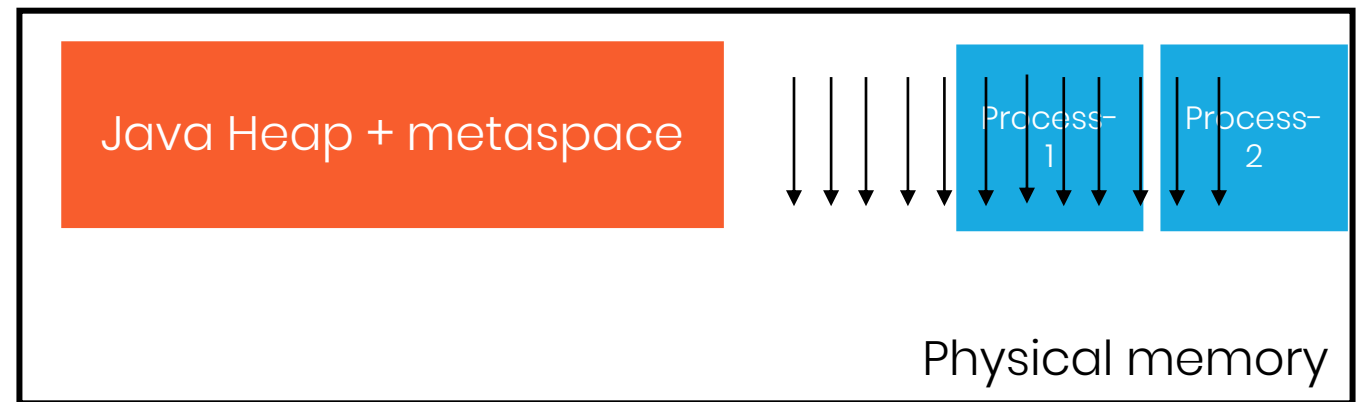


Key: Threads are created **outside heap, metaspace**

Solution:

1. Fix thread leak
2. Increase the Thread Limits Set at Operating System(`ulimit -u`)
3. Reduce Java Heap Size
4. Kills other processes
5. Increase physical memory size
6. Reduce thread stack size (`-Xss`).

Note: can cause `StackOverflowError`



8 types - OutOfMemoryError

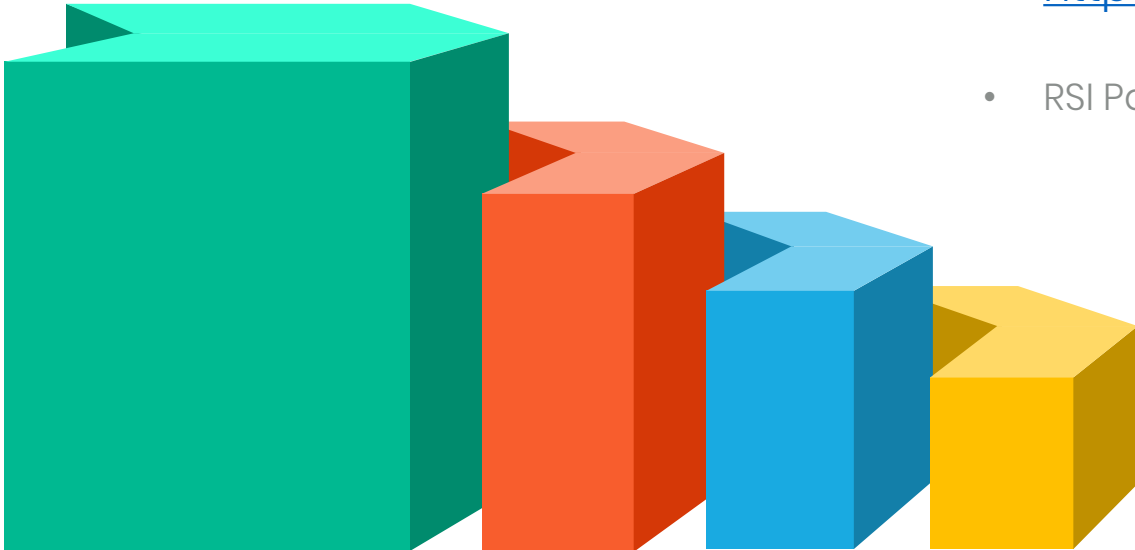
<https://blog.gceasy.io/2015/09/25/outofmemoryerror-beautiful-1-page-document/>

`java.lang.OutOfMemoryError: <type>`

- | | | | |
|----|--------------------------------------|----|---------------------------------------|
| 01 | Java heap space | 05 | Metaspace |
| 02 | GC overhead limit exceeded | 06 | Unable to create new native thread |
| 03 | Requested array size exceed VM limit | 07 | Kill process or sacrifice child |
| 04 | Permgen space | 08 | reason stack_trace_with_native method |

Troubleshooting unresponsive app

- <https://tinyurl.com/yywdmvyv>
- RSI Pattern – Same pattern, different problem.



Thread dump analysis Patterns

<https://blog.fastthread.io/category/thread-dump-patterns/>



RSI Pattern



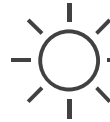
Treadmill Pattern



Leprechaun Pattern



All Roads leads to
Rome Pattern



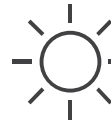
Traffic Jam Pattern



Athlete Pattern



Atherosclerosis Pattern

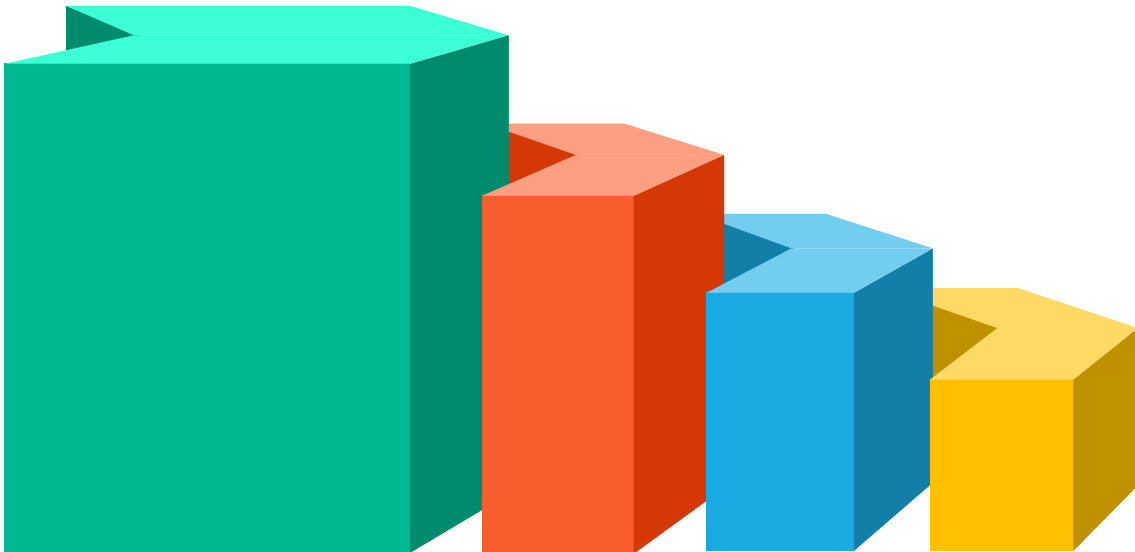


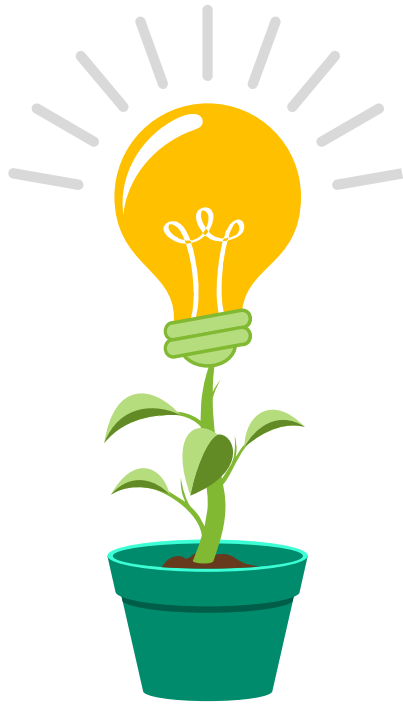
Several Scavengers Pattern



few more ...

Troubleshooting Memory problems





Enable GC Logs (always)

Till Java 8:

`-XX:+PrintGCDetails -XX:+PrintGCDateStamps -Xloggc:<file-path>`

From Java 9:

`-Xlog:gc*:file=<file-path>`

‘Free’ GC Log analysis tools

Freely available Garbage collection log analysis tools

01

GCeasy

<http://gceasy.io/>

02

GC Viewer

<https://github.com/chewiebug/GCViewer>

03

IBM GC & Memory visualizer

<https://developer.ibm.com/javasdk/tools/>

04

HP Jmeter

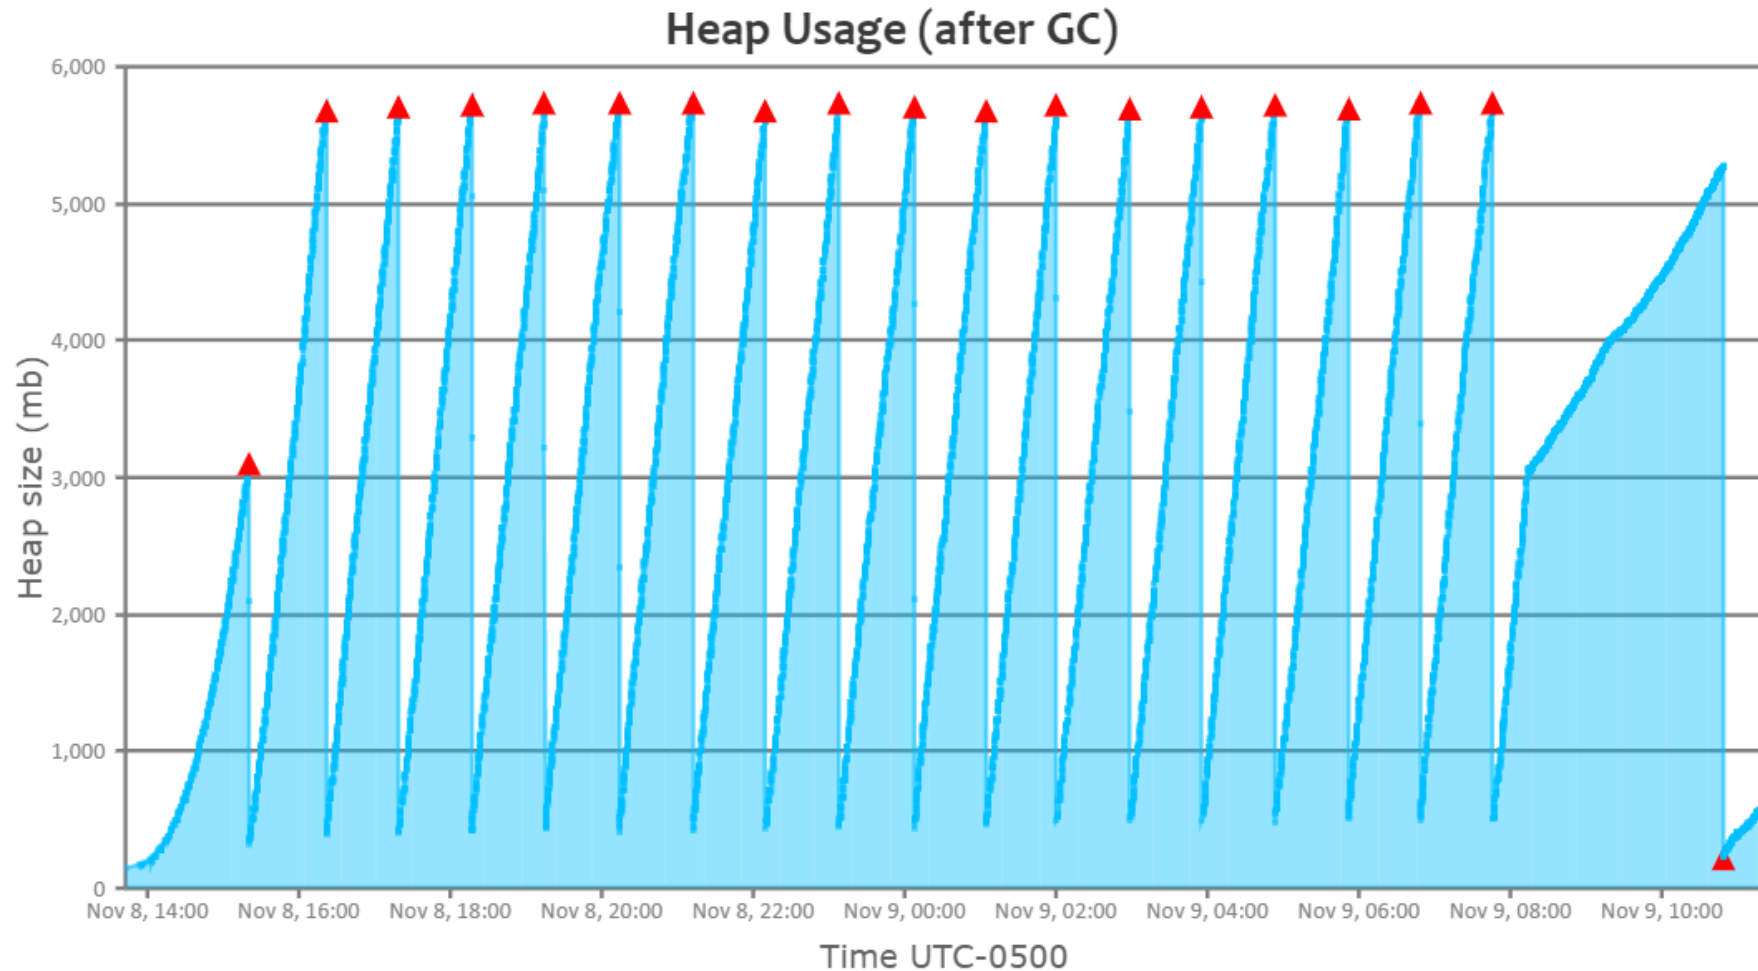
<https://h20392.www2.hp.com/portal/swdepot/displayProductInfo.do?productNumber=HPJMETER>

05

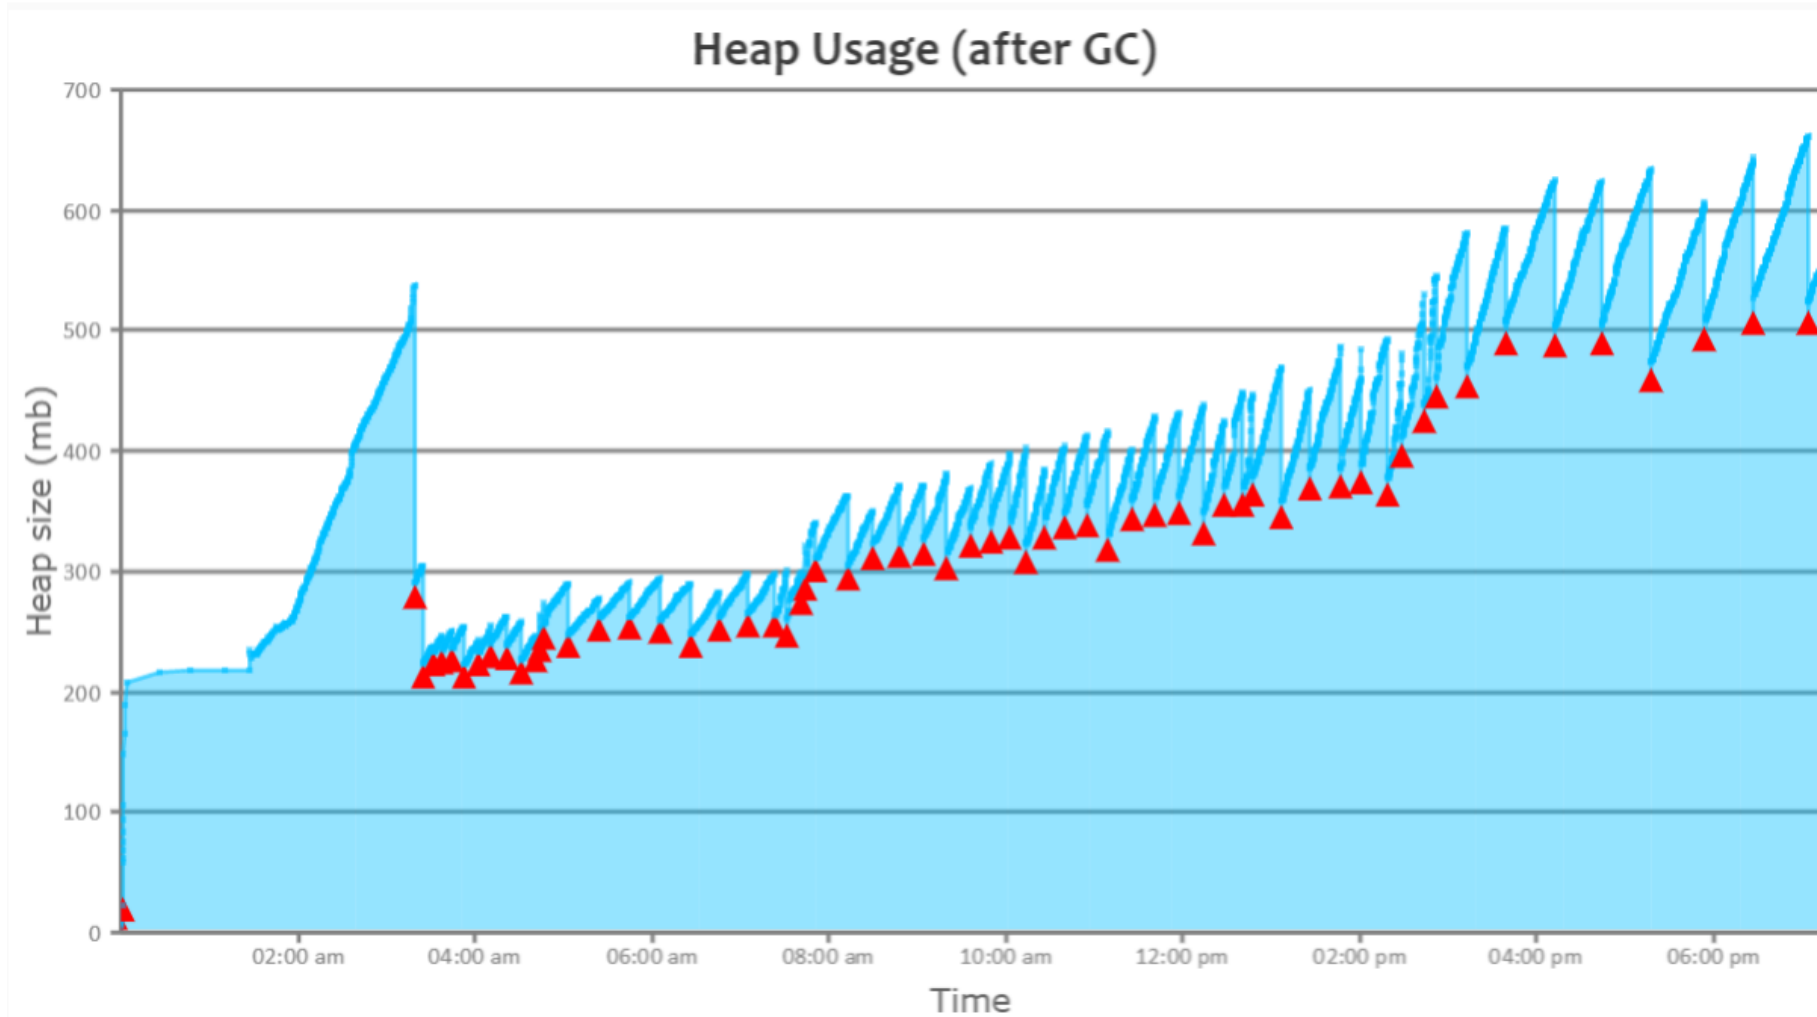
Google Garbage cat (cms)

<https://code.google.com/archive/a/eclipselabs.org/p/garbagecat>

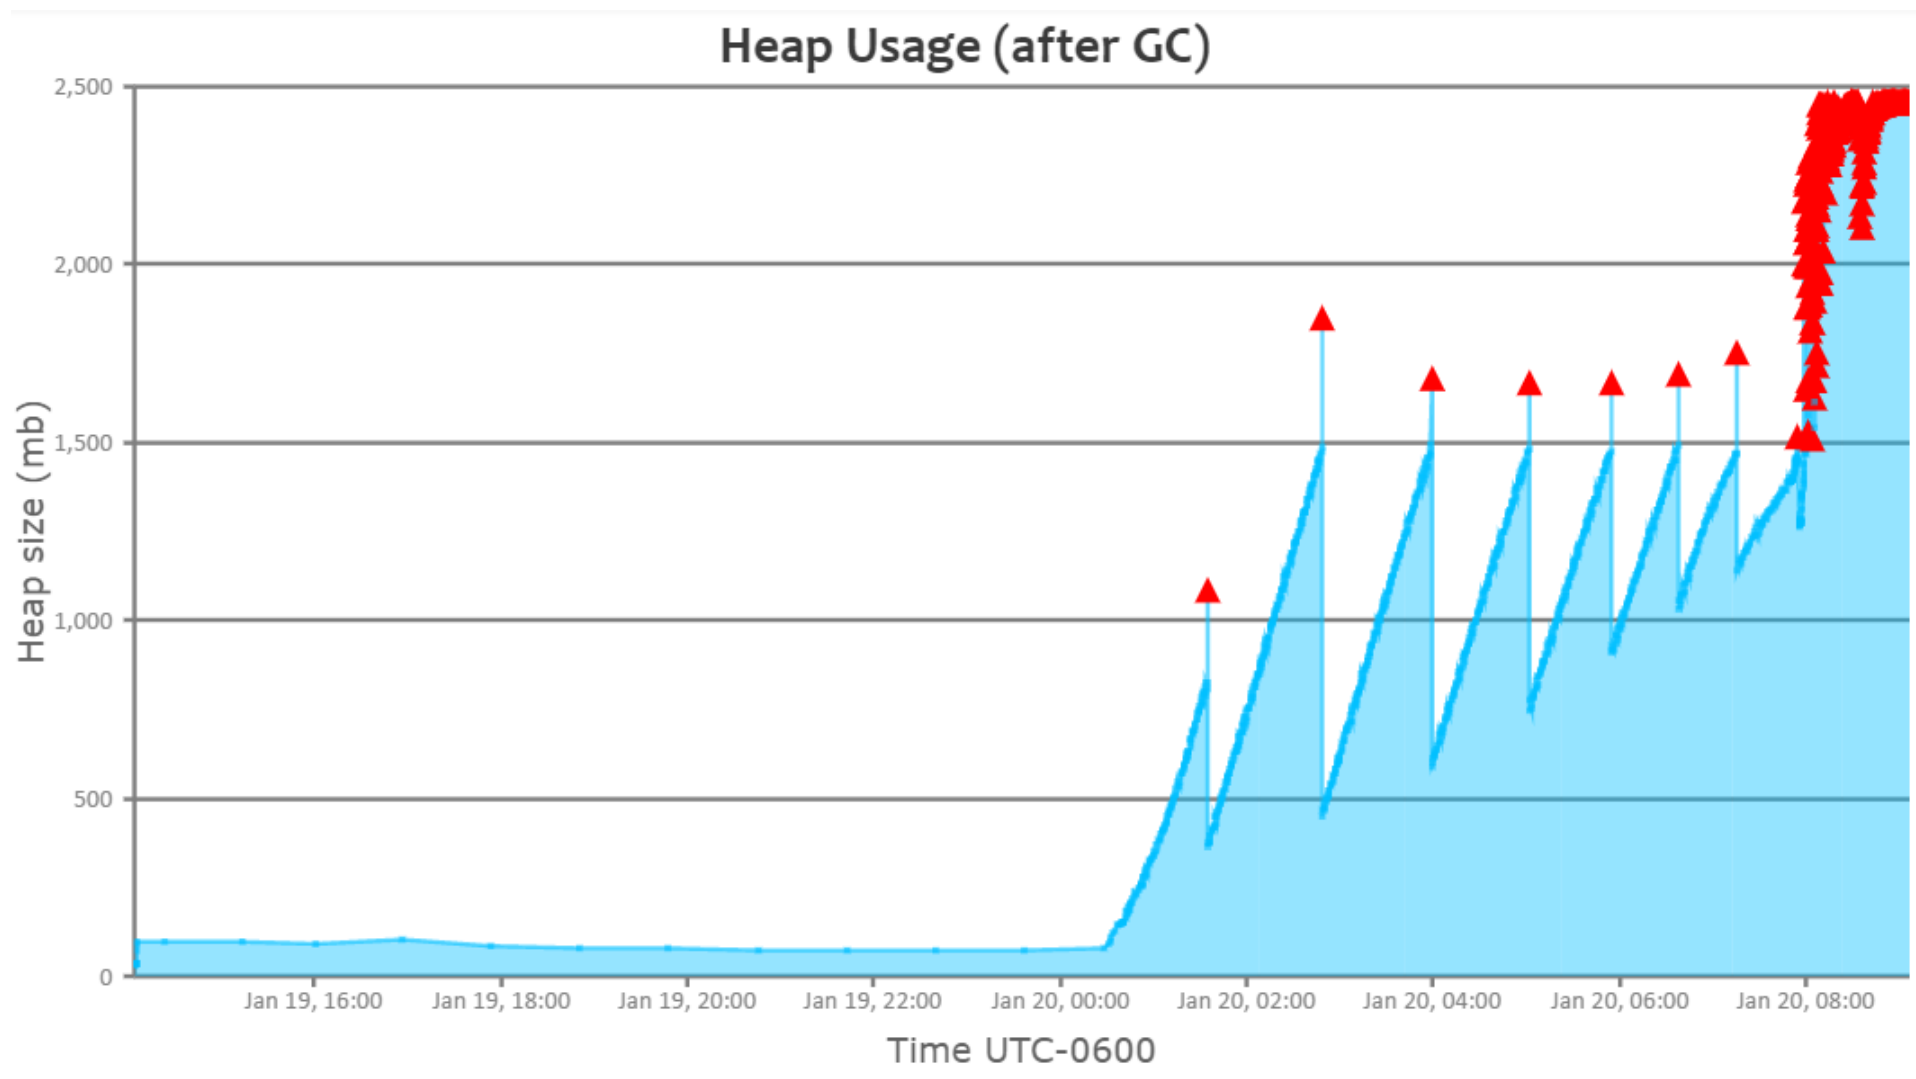
Heap usage graph



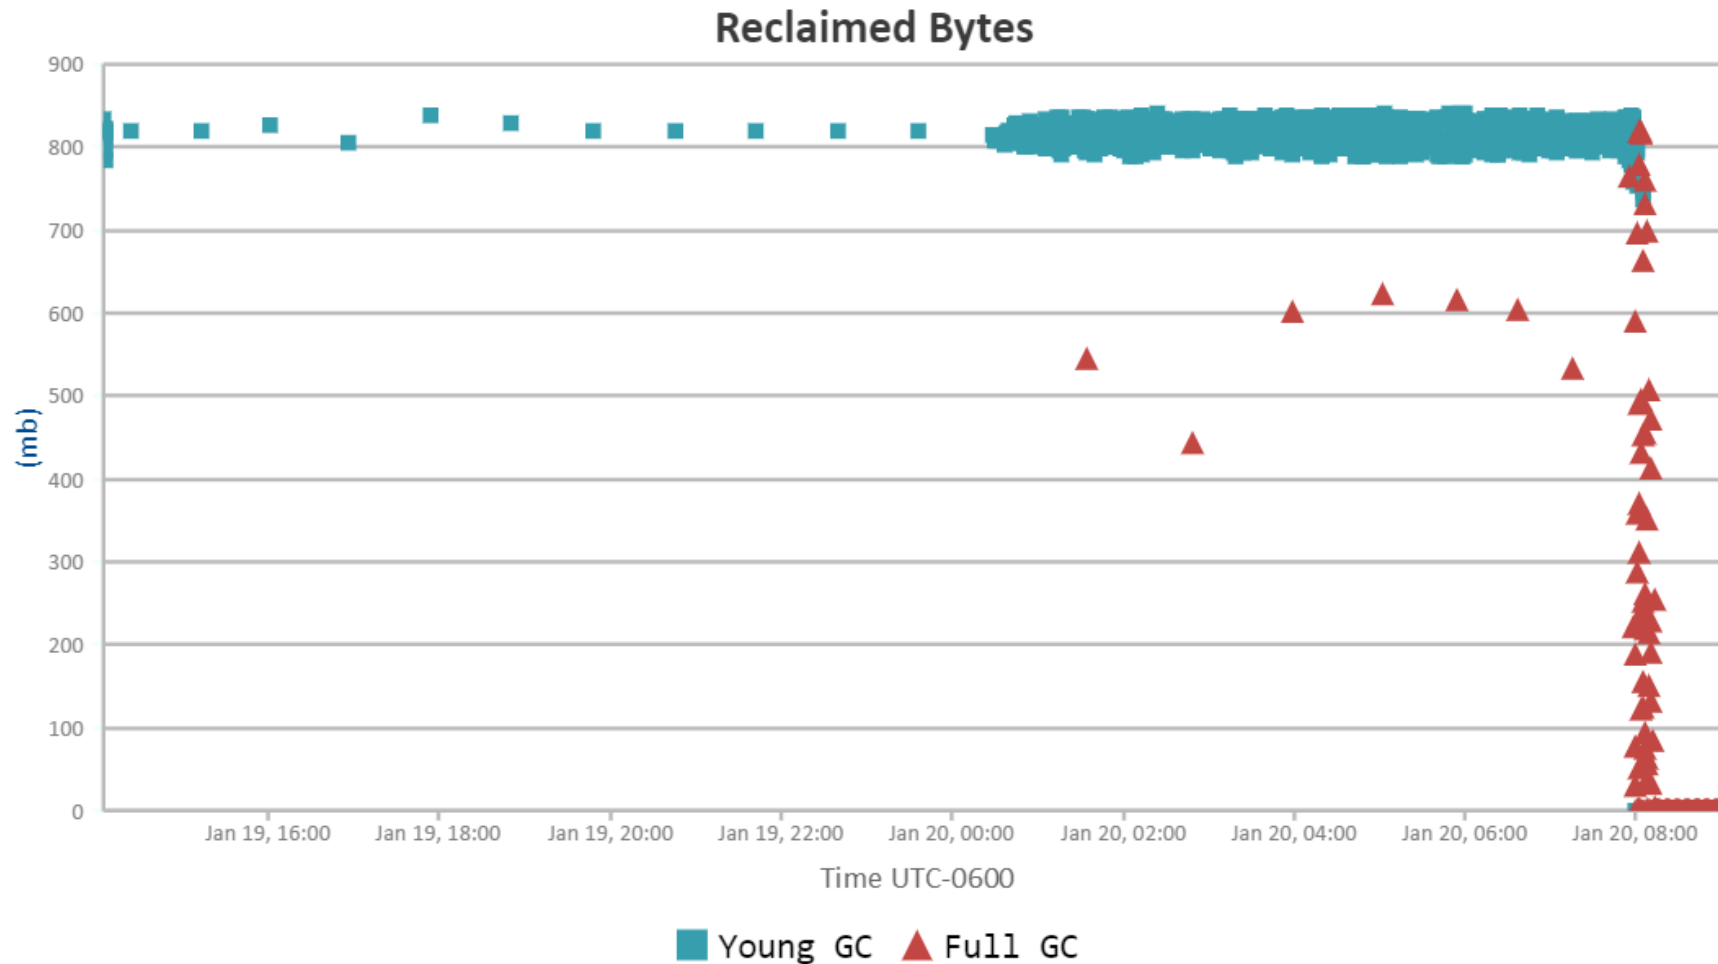
What is your observation?



Memory Problem



Corresponding – Reclaimed bytes chart



How to diagnose memory leak?



Capture heap dumps

`jmap -dump:live,file=<file-path> <pid>`

Example: `jmap -dump:live,file=/opt/tmp/AddressBook-heapdump.bin 37320`

`-XX:+HeapDumpOnOutOfMemoryError -XX:HeapDumpPath=/logs/heapdump`



Eclipse MAT, HeapHero

Two good tools to analyze memory leaks

Capture heap dumps

<https://blog.fastthread.io/2016/06/06/how-to-take-thread-dumps-7-options/>

01

jmap (since Java 5)

jmap -dump:live,file=<file-path> <pid>

02 ★

HeapDumpOnOutOfMemoryError

-XX:+HeapDumpOnOutOfMemoryError
-XX:HeapDumpPath=<file-path>

03

jVisualVM

JDK tool. Now Open source.
GUI based option.

04

IBM administrative console

If you are using WAS, this
option can be used

05

ThreadMXBean

Programmatic way to capture
thread dumps

06

APM Tools

Few APM Tools does provide
this support

07 ★

Jcmd (since Java 7)

jcmd <pid> GC.heap_dump
<file-path>



Micro-metrics

<https://blog.gceasy.io/2019/03/13/micrometrics-to-forecast-application-performance/>

Macro-Metrics

Can't forecast scalability, availability, performance problems



CPU

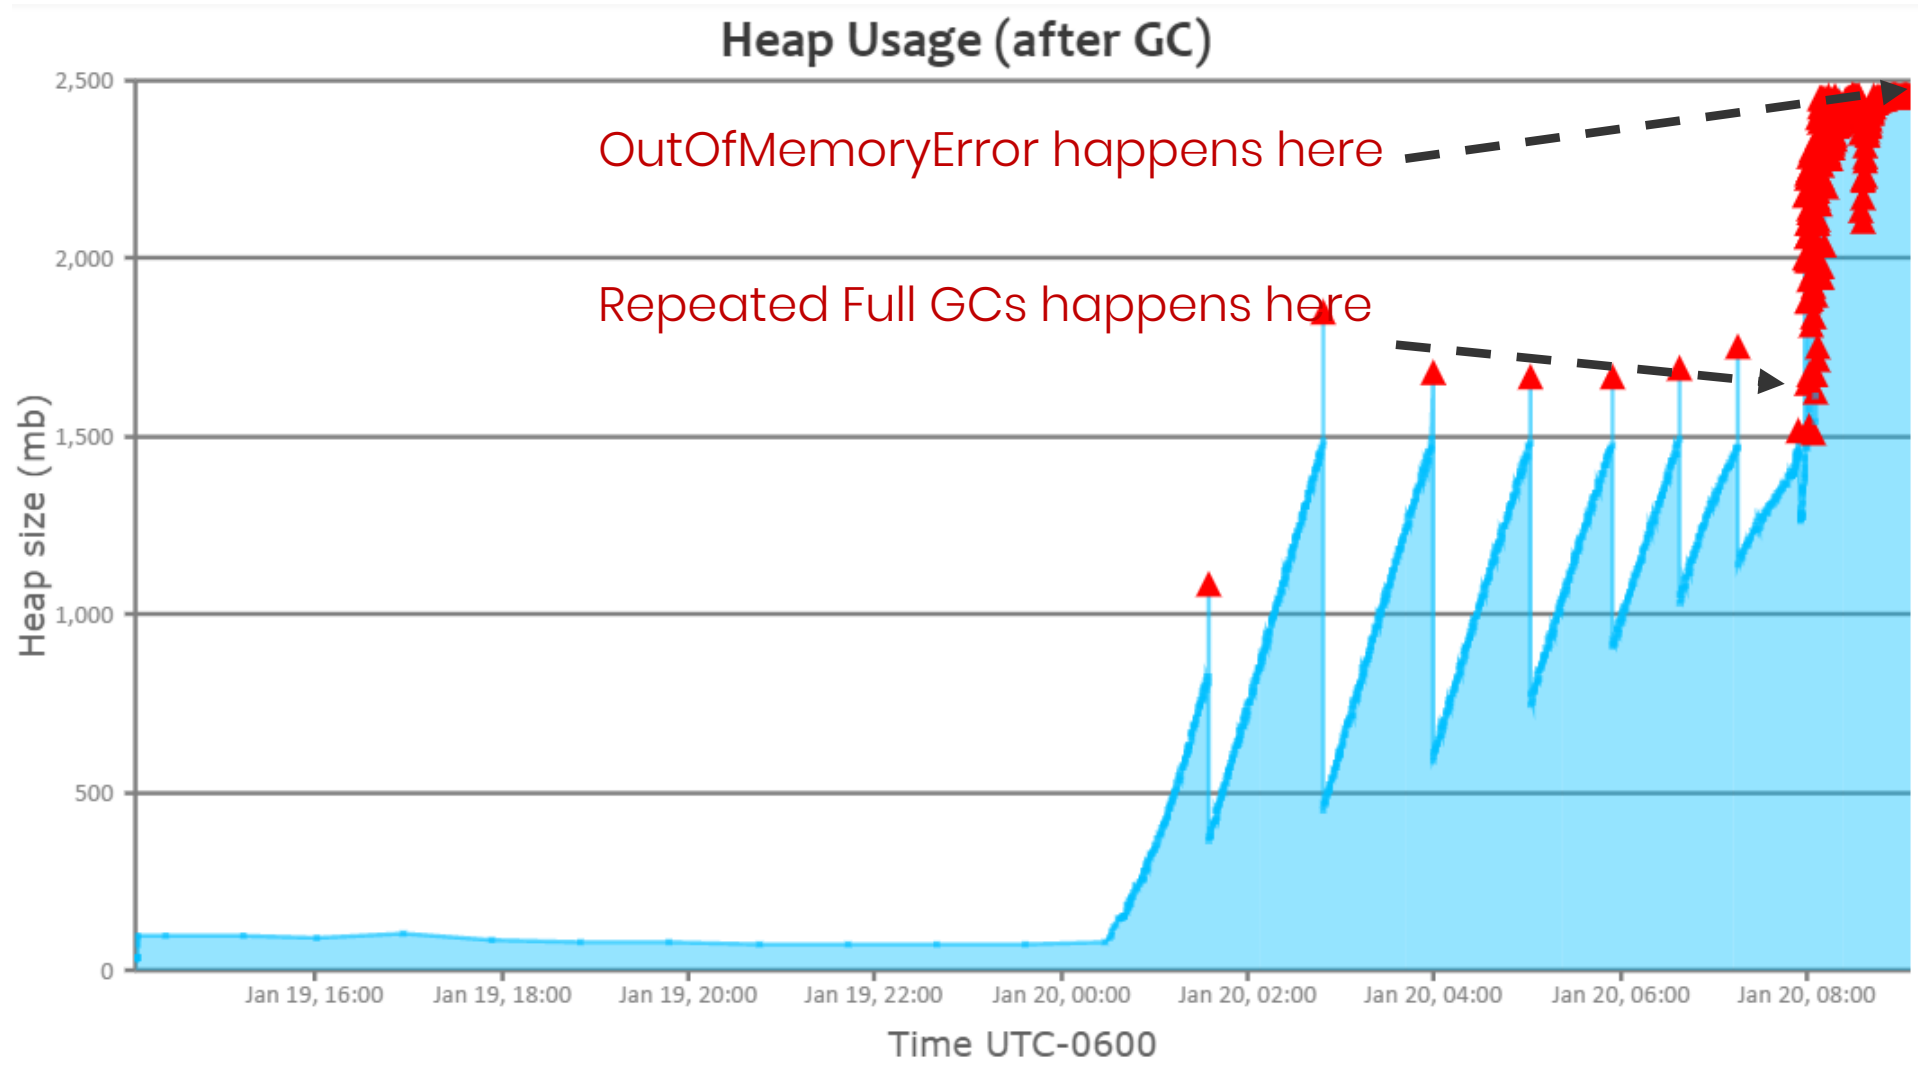


MEMORY



RESPONSE
TIME

Micro-metrics: Early Indicators



What are Micrometrics?

<https://blog.gceasy.io/2019/03/13/micrometrics-to-forecast-application-performance/>

GC Throughput

Amount time application spends in processing customer transactions vs amount of time application spend in doing GC

GC Latency

If pause time starts to increase, then it's an indication that app is suffering from memory problems

Object Reclamation rate

If number of objects created in unit time



File Descriptors

File descriptor is a handle to access: File, Pipe, Network Connections. If count grows it's a lead indicator that application isn't closing resources properly.

Thread States

If BLOCKED thread state count grows, it's an early indication that your application has potential to become unresponsive

Few more...

TCP/IP States, Hosts count, IOPS, ..



right data @ right time

What data to capture?



GC Log



Thread Dumps



Heap Dumps



top -H



top



netstat



vmstat



dmesg



ps



Disk Usage

IBM Script: <https://map.tinyurl.com/y4gz6o7q>
Captures all of the above artifacts

Thank you my friends!



Ram Lakshmanan



ram@tier1app.com



@tier1app



<https://www.linkedin.com/company/gceasy>